

## Chapter 2: Computer-System Structures (6<sup>th</sup> Edition)

肖 卿 俊

办公室：计算机楼532室

电邮：csqjxiao@seu.edu.cn


主页：<http://cse.seu.edu.cn/PersonalPage/csqjxiao>

电话：025-52091023




# Chapter 2: Computer-System Structures (6<sup>th</sup> Ed.)

- Computer System Operation
- I/O Structure
- Storage Structure
- Storage Hierarchy (层次结构)
- Hardware Protection
- Network Structure





# Computer-System Architecture




- How does the OS get into system?
- How do we request for services?
- How does the OS know something has happened





# Computer-System Operation


- I/O devices and the CPU can execute concurrently.
- Each device controller is in charge of a particular device type.
- Each device controller has a local buffer.





# Computer-System Operation (Cont.)

- CPU moves data from/to main memory to/from local buffers of controllers
- I/O is from the device to local buffer of controller.
- Device controller informs CPU that it has finished its operation by causing an *interrupt*.






# Interrupt

- An event that requires the attention of the OS is an *interrupt*. These events include the completion of an I/O, a keypress, a request for service, a division by zero and so on.
- Interrupts may be triggered by software.
  - ◆ An interrupt generated by software (*i.e.*, division by 0, page fault, debug breakpoint) is usually referred to as a *trap*.
- Modern operating systems are *interrupt driven*, meaning the OS is in action only if an interrupt occurs.






# What is Interrupt driven?





- The OS is activated by an interrupt.
- The executing program is suspended.
- Control is transferred to the OS.
- Program continues when the service completes





# Interrupt Time Line For a Single Process Doing Output








# Common Functions of Interrupts

- Interrupt transfers control to the interrupt service routine generally, through the *interrupt vector*, which contains the addresses of all the service routines.
- Interrupt architecture must save the address of the interrupted instruction.
- Incoming interrupts are *disabled* while another interrupt is being processed to prevent a *lost interrupt*.






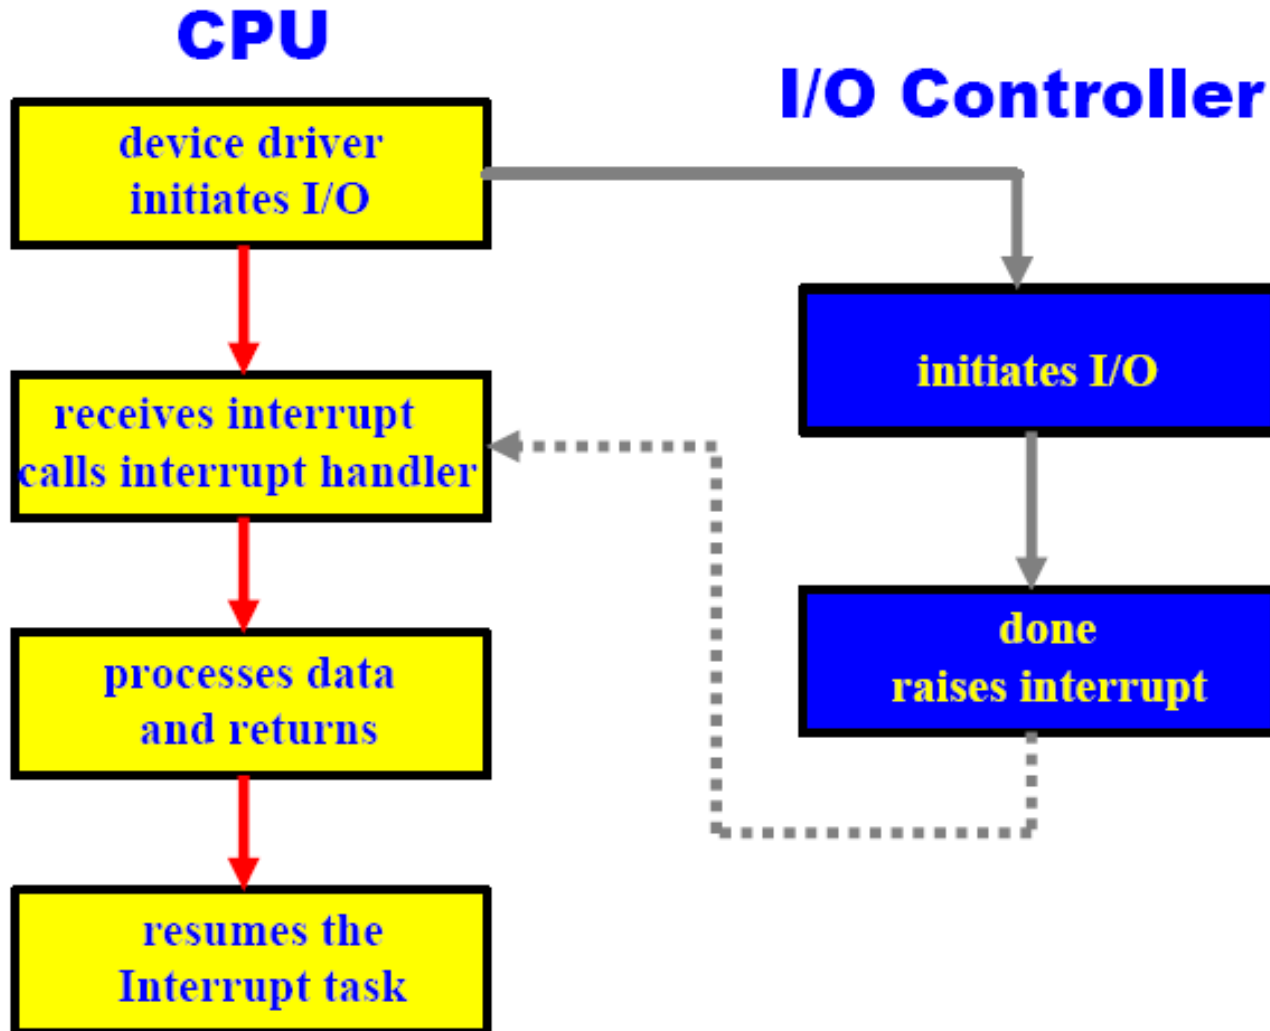
# Interrupt Handling

- The operating system preserves the state of the CPU by storing registers and the program counter.
- Determines which type of interrupt has occurred:
  - ◆ *polling*
  - ◆ *vectored* interrupt system
- Separate segments of code determine what action should be taken for each type of interrupt






# I/O Interrupt




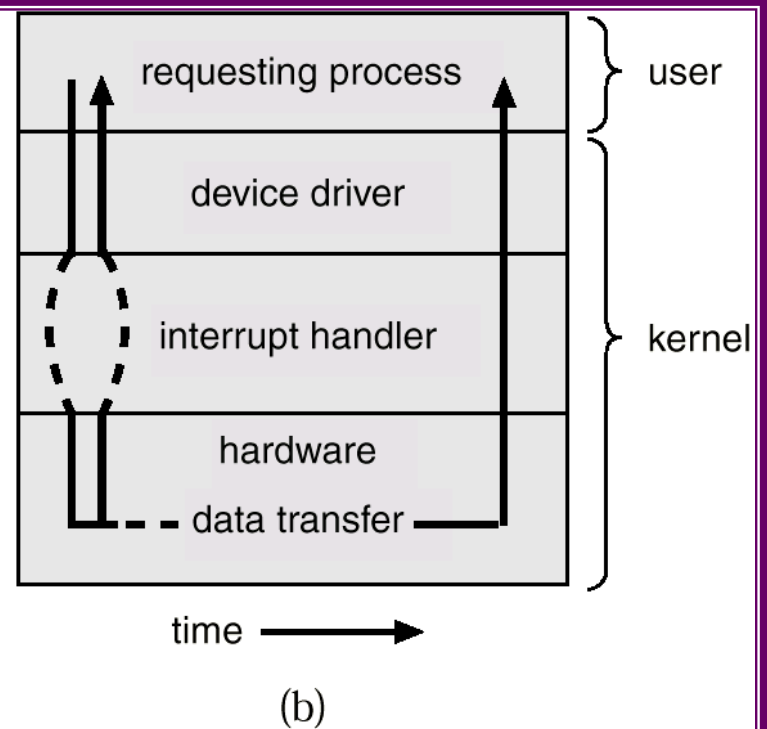


# Two I/O Methods: Synchronous vs. Asynchronous

Synchronous




Asynchronous





# Synchronous I/O from User View

- After I/O starts, control returns to user program only upon I/O completion.
  - ◆ Wait instruction idles the user process until the next interrupt
  - ◆ Wait loop (contention for memory access).
  - ◆ At most one I/O request is outstanding at a time, no simultaneous I/O processing.






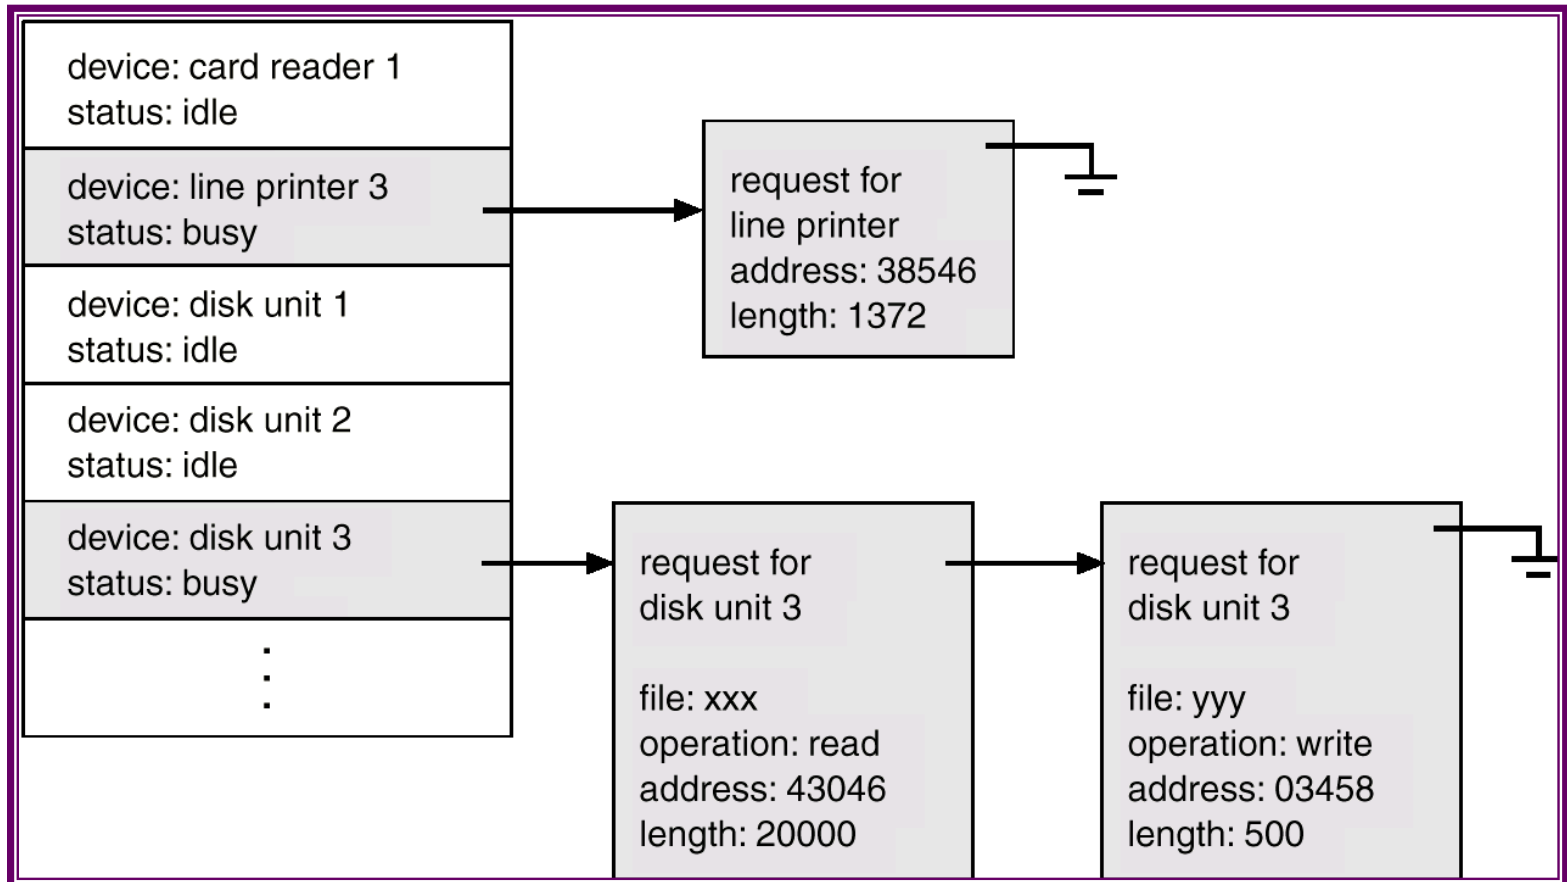
# Asynchronous I/O from User View

- After I/O starts, control returns to user program without waiting for I/O completion.
  - ◆ *System call* – request to the operating system to allow user to wait for I/O completion.
  - ◆ *Device-status table* contains entry for each I/O device indicating its type, address, and state.
  - ◆ Operating system indexes into I/O device table to determine device status and to modify table entry to include interrupt.
- How does the user program receive the notification of I/O completion?





# Device-Status Table





# Direct Memory Access Structure (1/2)

- Used for high-speed I/O devices able to transmit information at close to memory speeds.
- Device controller transfers blocks of data from buffer storage directly to main memory without CPU intervention.
- Only one interrupt is generated per block, rather than the one interrupt per byte.

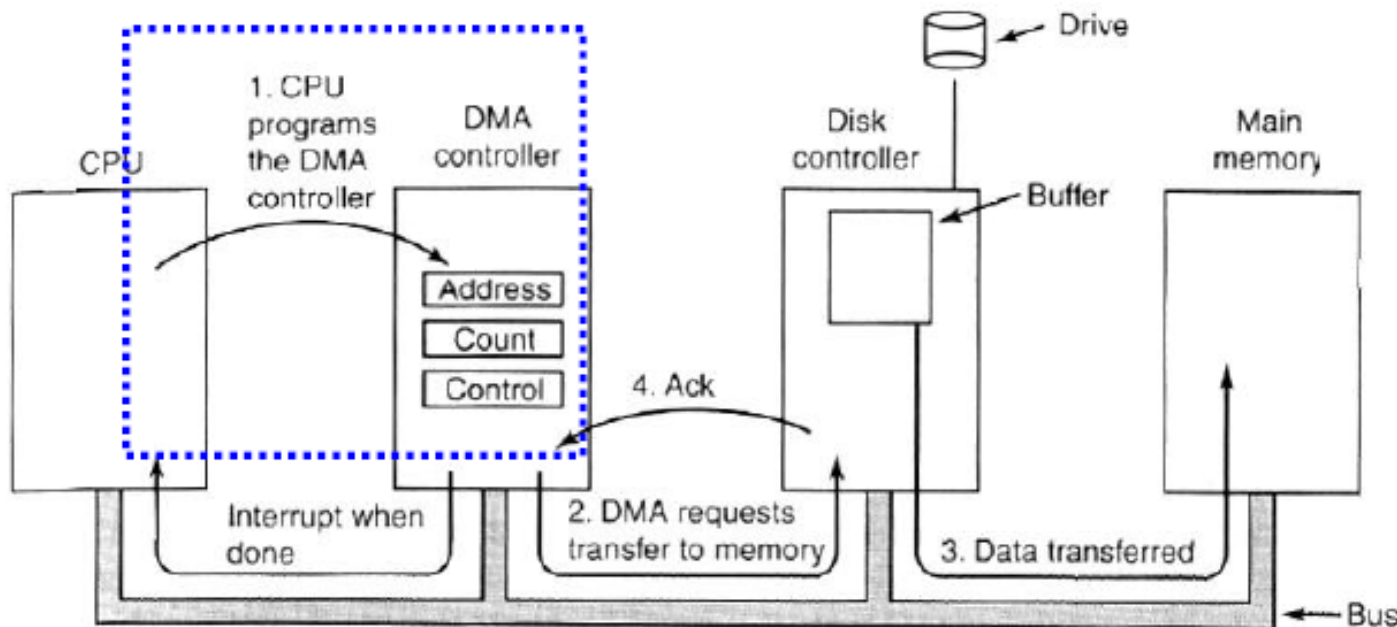






# Direct Memory Access Structure (2/2)

The CPU gives the controller (1) **disk address**, (2) **memory address** for storing the block, and (3) a **byte count**. Then, the CPU goes back to work.







# Storage Structure

- Main memory – only large storage media that the CPU can access directly.
- Secondary storage – extension of main memory that provides large nonvolatile storage capacity.
- Magnetic disks – rigid metal or glass platters covered with magnetic recording material
  - ◆ Disk surface is logically divided into *tracks*, which are subdivided into *sectors*.
  - ◆ The *disk controller* determines the logical interaction between the device and the computer.





# Moving-Head Disk Mechanism






# Storage Hierarchy

- Storage systems organized in hierarchy.
  - ◆ Speed
  - ◆ Cost
  - ◆ Volatility
- *Caching* – copying information into faster storage system; main memory can be viewed as a last *cache* for secondary storage.






# Storage-Device Hierarchy



Speed ↑


Volume ↓






# Caching

- Use of high-speed memory to hold recently-accessed data.
- Requires a *cache management* policy.
- Caching introduces another level in storage hierarchy. This requires data that is simultaneously stored in more than one level to be *consistent*.





# Migration of A From Disk to Register



How to maintain the Cache Consistency?





# Hardware Protection

- Dual-Mode Operation
- I/O Protection
- Memory Protection
- CPU Protection







# Dual-Mode Operation


- Sharing system resources requires operating system to ensure that an incorrect program cannot cause other programs to execute incorrectly.
- Provide hardware support to differentiate between at least two modes of operations.
  1. *User mode* – execution done on behalf of a user. (用户态或目态)
  2. *Monitor mode* (also *kernel mode* or *system mode*) – execution done on behalf of operating system. (核心态或管态)







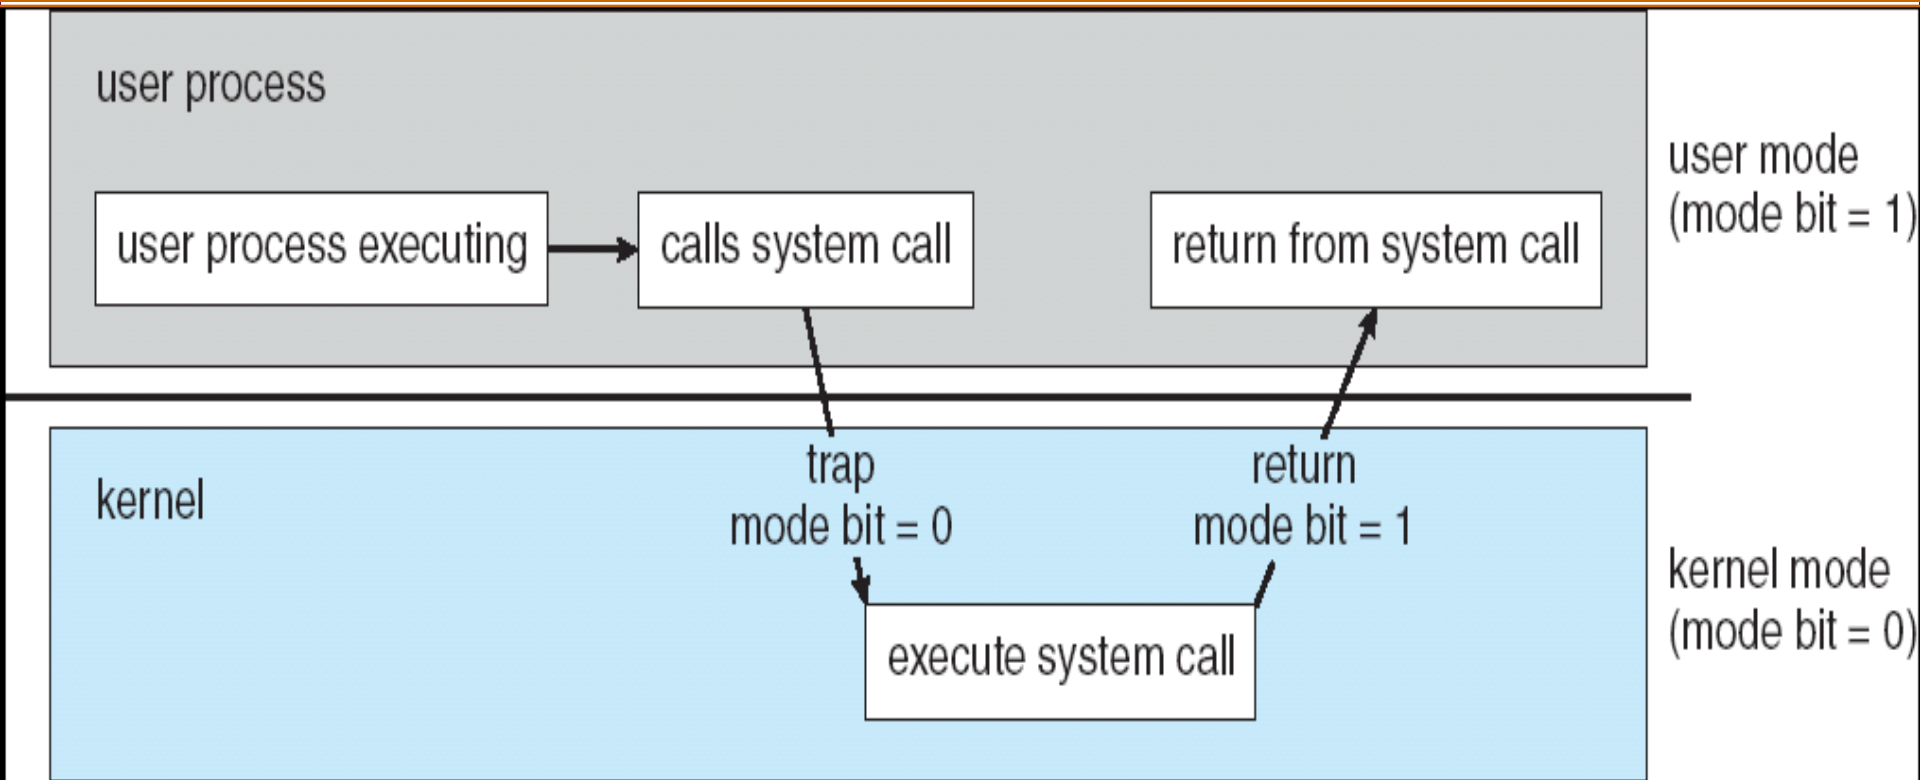
# Dual-Mode Operation (Cont.)

- *Mode bit* added to computer hardware to indicate the current mode: monitor (0) or user (1).
- When an interrupt or fault occurs hardware switches to monitor mode.



*Privileged instructions can be issued only in monitor mode.*










# I/O Protection

- All I/O instructions are privileged instructions (特权指令), e.g., send commands to IO controllers
- Must ensure that a user program could never gain control of the computer in monitor mode (i.e., a user program that, as part of its execution, stores a new address in the interrupt vector).





# Use of A System Call to Perform I/O






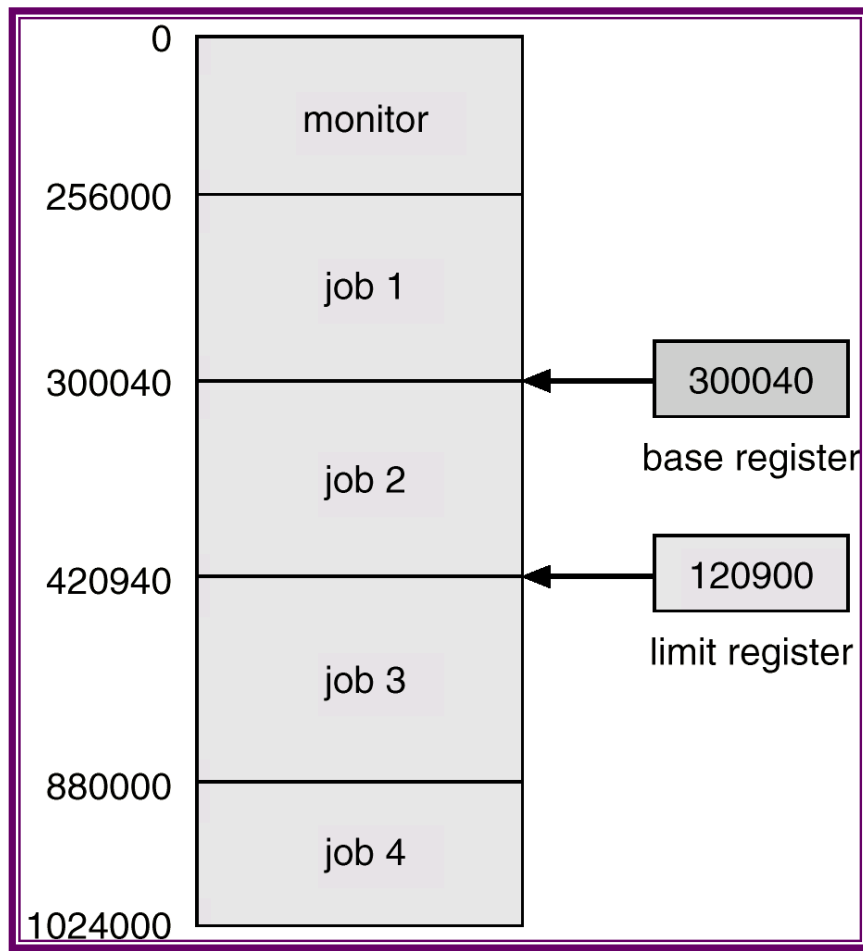
# Memory Protection

- Must provide memory protection at least for the interrupt vector and the interrupt service routines.
- In order to have memory protection, add two registers that determine the range of legal addresses a program may access:
  - ◆ **Base register** (基址寄存器) – holds the smallest legal physical memory address.
  - ◆ **Limit register** (界限寄存器) – contains the size of the range
- Memory outside the defined range is protected.




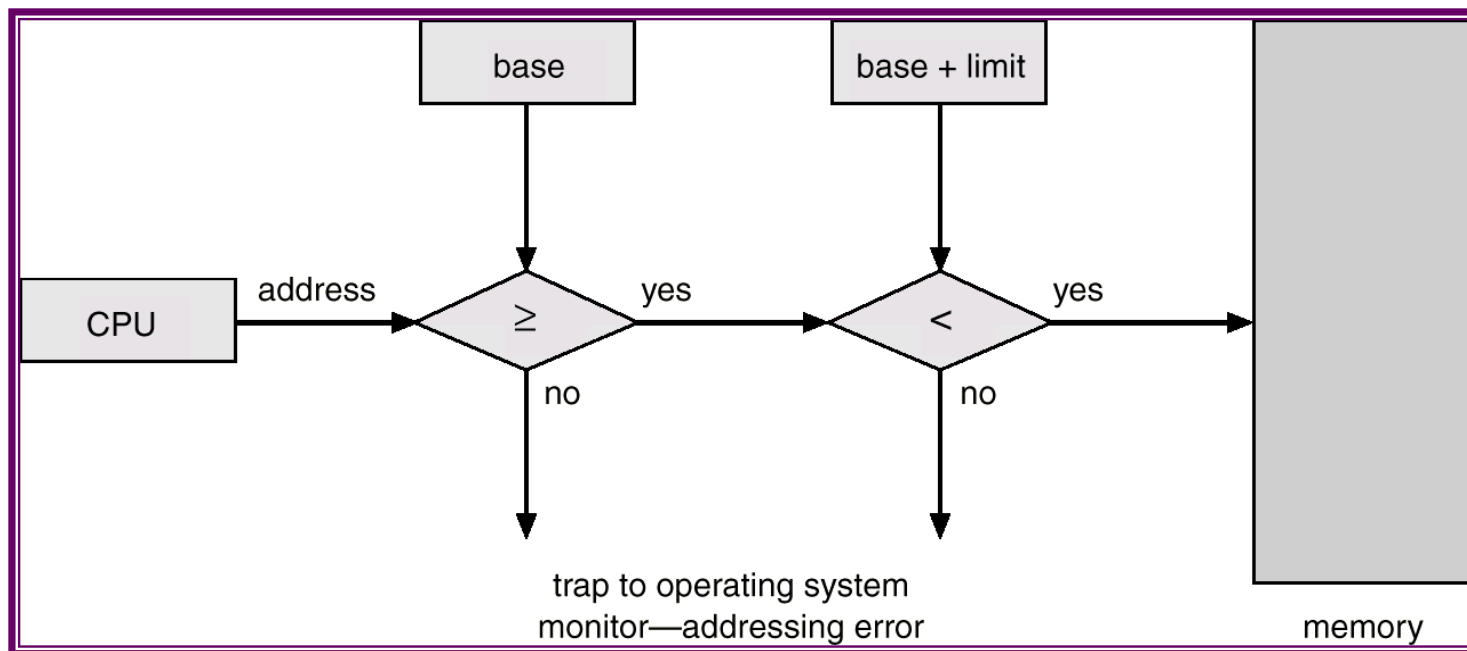


# Use of A Base and Limit Register





# Hardware Address Protection







# Hardware Protection

- When executing in monitor mode, the operating system has unrestricted access to both monitor and user's memory.
- The load instructions for the *base* and *limit* registers are privileged instructions (特权指令).





# CPU Protection

- *Timer* – interrupts computer after specified period to ensure operating system maintains control.
  - ◆ Timer is decremented every clock tick.
  - ◆ When timer reaches the value 0, an interrupt occurs.
- Timer commonly used to implement time sharing.
- Timer also used to compute the current time.
- Load-timer is a privileged instruction.

