

## Chapter 6: Process Synchronization


肖卿俊

办公室：计算机楼532室

电邮：csqjxiao@seu.edu.cn


主页：<http://cse.seu.edu.cn/PersonalPage/csqjxiao>


电话：025-52091023



# Chapter 6: Process Synchronization

- Background
- The Critical-Section Problem
- Synchronization Hardware
- Semaphores
- Classical Problems of Synchronization
- Monitors
- Synchronization Examples







# Background


- Concurrent access to shared data may result in data inconsistency.
- Maintaining **data consistency** requires mechanisms to ensure the **orderly execution** of cooperating processes.
- Synchronization primitive for shared memory programming --- Lock

◆ Exclusive Lock




◆ Shared Lock






## Background (cont.)

- Other synchronization primitive for shared memory programming --- Barrier



- As for this course, we focus primarily on the lock primitive
- Ideas are easy. Implementation is hard.





# Revisit the Shared-memory Producer Consumer Problem

- Shared-memory solution to bounded-buffer problem (Chapter 3) allows at most  $n - 1$  items in buffer at the same time. A solution, where all  $N$  buffers are used is not simple.


## Producer:


```
while (1) {  
    while (((in+1) % BUFFER_SIZE) == out) ;  
    .....  
    in = (in+1) % BUFFER_SIZE;  
}
```

## Consumer:

```
while (1) {  
    while (in == out) ;  
    .....  
    out = (out+1) % BUFFER_SIZE;  
}
```

- ◆ Suppose that we modify the producer-consumer code by adding a variable *counter*







# Bounded-Buffer Solution

## ■ Shared data

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;
item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
int counter = 0;
```







# Bounded-Buffer Solution

## ■ Producer process

```
item nextProduced;
```

```
while (1) {  
    while (counter == BUFFER_SIZE)  
        ; /* do nothing */  
    buffer[in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```






# Bounded-Buffer Solution


## ■ Consumer process

```
item nextConsumed;
```

```
while (1) {  
    while (counter == 0)  
        ; /* do nothing */  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
}
```








## Critical Shred Data


- Counter is a piece of critical shared data
- The statements

**counter++;**  
**counter--;**

must be performed *atomically*.

- Atomic operation means an operation that completes in its entirety without interruption.





# Difficult to Implement the Atomic Guarantee

- However, the statement “**count++**” may be implemented in machine language as:

**register1 = counter**

**register1 = register1 + 1**


**counter = register1**


- The statement “**count--**” may be implemented in machine language as:

**register2 = counter**

**register2 = register2 – 1**


**counter = register2**






# Potential Data Inconsistency

- If both the producer and consumer attempt to update the buffer concurrently, the assembly language statements may get interleaved.
- Interleaving depends upon how the producer and consumer processes are scheduled.





# Potential Data Inconsistency

- Assume **counter** is initially 5. One interleaving of statements is:

producer: **register1 = counter** (*register1 = 5*)

producer: **register1 = register1 + 1** (*register1 = 6*)

consumer: **register2 = counter** (*register2 = 5*)


consumer: **register2 = register2 - 1** (*register2 = 4*)

producer: **counter = register1** (*counter = 6*)

consumer: **counter = register2** (*counter = 4*)

- The value of **count** may be either 4 or 6, where the correct result should be 5.





Producer

**register1 = counter**

**register1 = register1 + 1**


**counter = register1**


Consumer

**register2 = counter**

**register2 = register2 - 1**


**counter = register2**






# Concept of Race Condition


- **Race condition** occurs, if:
  - ◆ **two or more** processes/threads access and manipulate the **same** data **concurrently**, and
  - ◆ the outcome of the execution **depends on the particular order** in which the access takes place.
  
- To prevent race conditions, concurrent processes must be **synchronized**.






# Chapter 6: Process Synchronization


- Background
- The Critical-Section Problem
- Synchronization Hardware
- Semaphores
- Classical Problems of Synchronization
- Monitors
- Synchronization Examples






# The Critical-Section Problem

- Multiple processes all competing to use some shared data
- Each process has a code segment, called *critical section*, in which the shared data is accessed.
- Problem – ensure that when one process is executing in its critical section, no other process is allowed to execute in its critical section.









# Critical Section and Mutual Exclusion

- Thus, the execution of critical sections must be *mutually exclusive* (e.g., at most one process can be in its critical section at any time).
- The *critical-section problem* is to design a protocol that processes can use to cooperate.





# The Critical Section Protocol

```
do {
```



entry section


critical section


exit section



```
} while (1);
```


- A critical section protocol consists of **two** parts: an **entry section (or lock)** and an **exit section (or unlock)**.
- Between them is the critical section that must run in a **mutually exclusive** way.






# Solution to Critical-Section Problem

- Any solution to the critical section problem must satisfy the following three conditions:
  - ◆ **Mutual Exclusion**
  - ◆ **Progress**
  - ◆ **Bounded Waiting**
  
- Moreover, the solution cannot depend on **relative speed** of processes and **scheduling policy**.






# Mutual Exclusion


- If a process **P** is executing in its critical section, then **no** other processes can be executing in their critical sections.
- The **entry protocol** should be capable of blocking processes that wish to enter but cannot.
- Moreover, when the process that is executing in its critical section exits, the **entry protocol** must be able to know this fact and allows a waiting process to enter.






# Progress


- If **no** process is executing in its critical section and some processes wish to enter their critical sections, then
  - ◆ Only those processes that are waiting to enter can participate in the competition (to enter their critical sections).
  - ◆ No other process can influence this decision.
  - ◆ This decision cannot be postponed indefinitely.






# Bounded Waiting

- After a process made a request to enter its critical section and before it is granted the permission to enter, there exists a *bound* on the number of times that other processes are allowed to enter.
- Hence, even though a process may be blocked by other waiting processes, it will not be waiting forever.
  - Assume that each process executes at a nonzero speed
  - No assumption concerning relative speed of the  $n$  processes






# Initial Attempts to Solve Problem

- Consider a special case of only 2 processes,  $P_0$  and  $P_1$
- General structure of process  $P_i$  (the other process  $P_j$ )

```
do {  
    entry section  
    critical section  
    exit section  
    remainder section  
} while (1);
```

- Processes may share some common variables to synchronize their actions.





# Algorithm 1

- Shared variable:

- ◆ **boolean lock;**  
initially **lock = false**

- ◆ **lock = true**  $\Rightarrow$  the critical section has been locked

- Process  $P_i$ :

- do {    **while (lock) ;**  
          **lock = true;**  
          critical section  
          **lock = false;**  
          remainder section  
      } **while (1);**

- Does not satisfy **mutual exclusion**

<https://en.wikipedia.org/wiki/Test-and-set>







## Algorithm 2

### ■ Shared variables

◆ **boolean flag[2];**

initially **flag[0] = flag[1] = false.**

◆ **flag[i] = true**  $\Rightarrow P_i$  wants to enter its critical section


### ■ Process $P_i$

```
do {   flag[i] = true;
      while (flag[j]) ;
      critical section
      flag[i] = false;
      remainder section
```


```
} while (1);
```

### ■ Satisfies mutual exclusion, but not progress requirement





# Deadlock Problem of Algorithm 2







# Is the Following Algorithm Correct?

- What if we change the location of the statement: **flag[i] = true**?
- Process  $P_i$ :

```
do {    while (flag[j]) ;  
    flag[i] = true;  
    critical section  
    flag[i] = false;  
    remainder section  
} while (1);
```
- Does not satisfy **mutual exclusion**





## Algorithm 3

■ Shared variables:

◆ **int victim;**      initially **victim = i** (or **victim = j**)

■ Process  $P_i$ :

```
do {      victim = i;  
        while (victim == i) ;  
        critical section  
        // do nothing for CS exit  
        remainder section  
    } while (1);
```

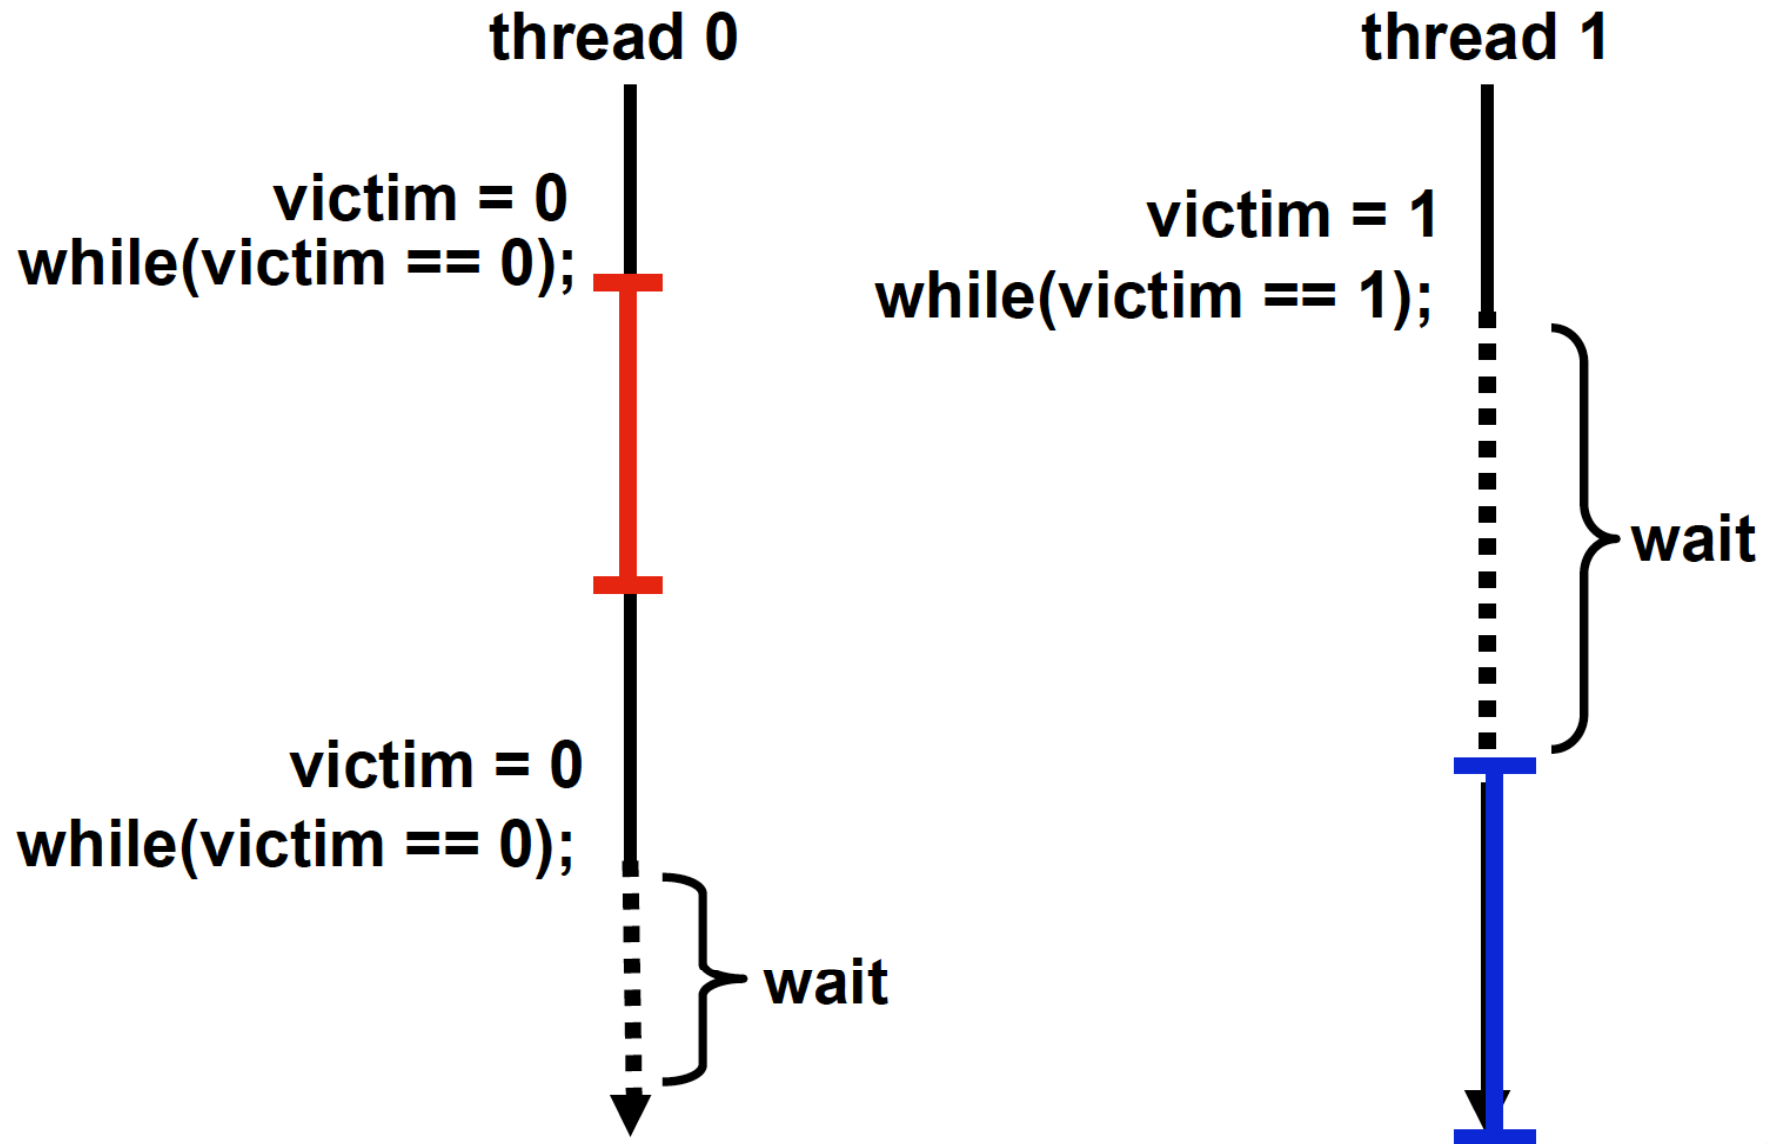
■ Processes are forced to run in an alternating way


■ Satisfies **mutual exclusion**, but not **progress**








# Alternating and Atomic Execution of Algorithm 3





# Deadlock of Algorithm 3





# Peterson's Algorithm

- Combined shared variables of algorithms 2, 3.


- Process  $P_i$


```
do {   flag[i] = true;   // I'm interested
      victim = i;       // you go first
      while (flag[j] and victim == i) ;
      critical section
      flag[i] = false;
      remainder section
```

```
} while (1);
```

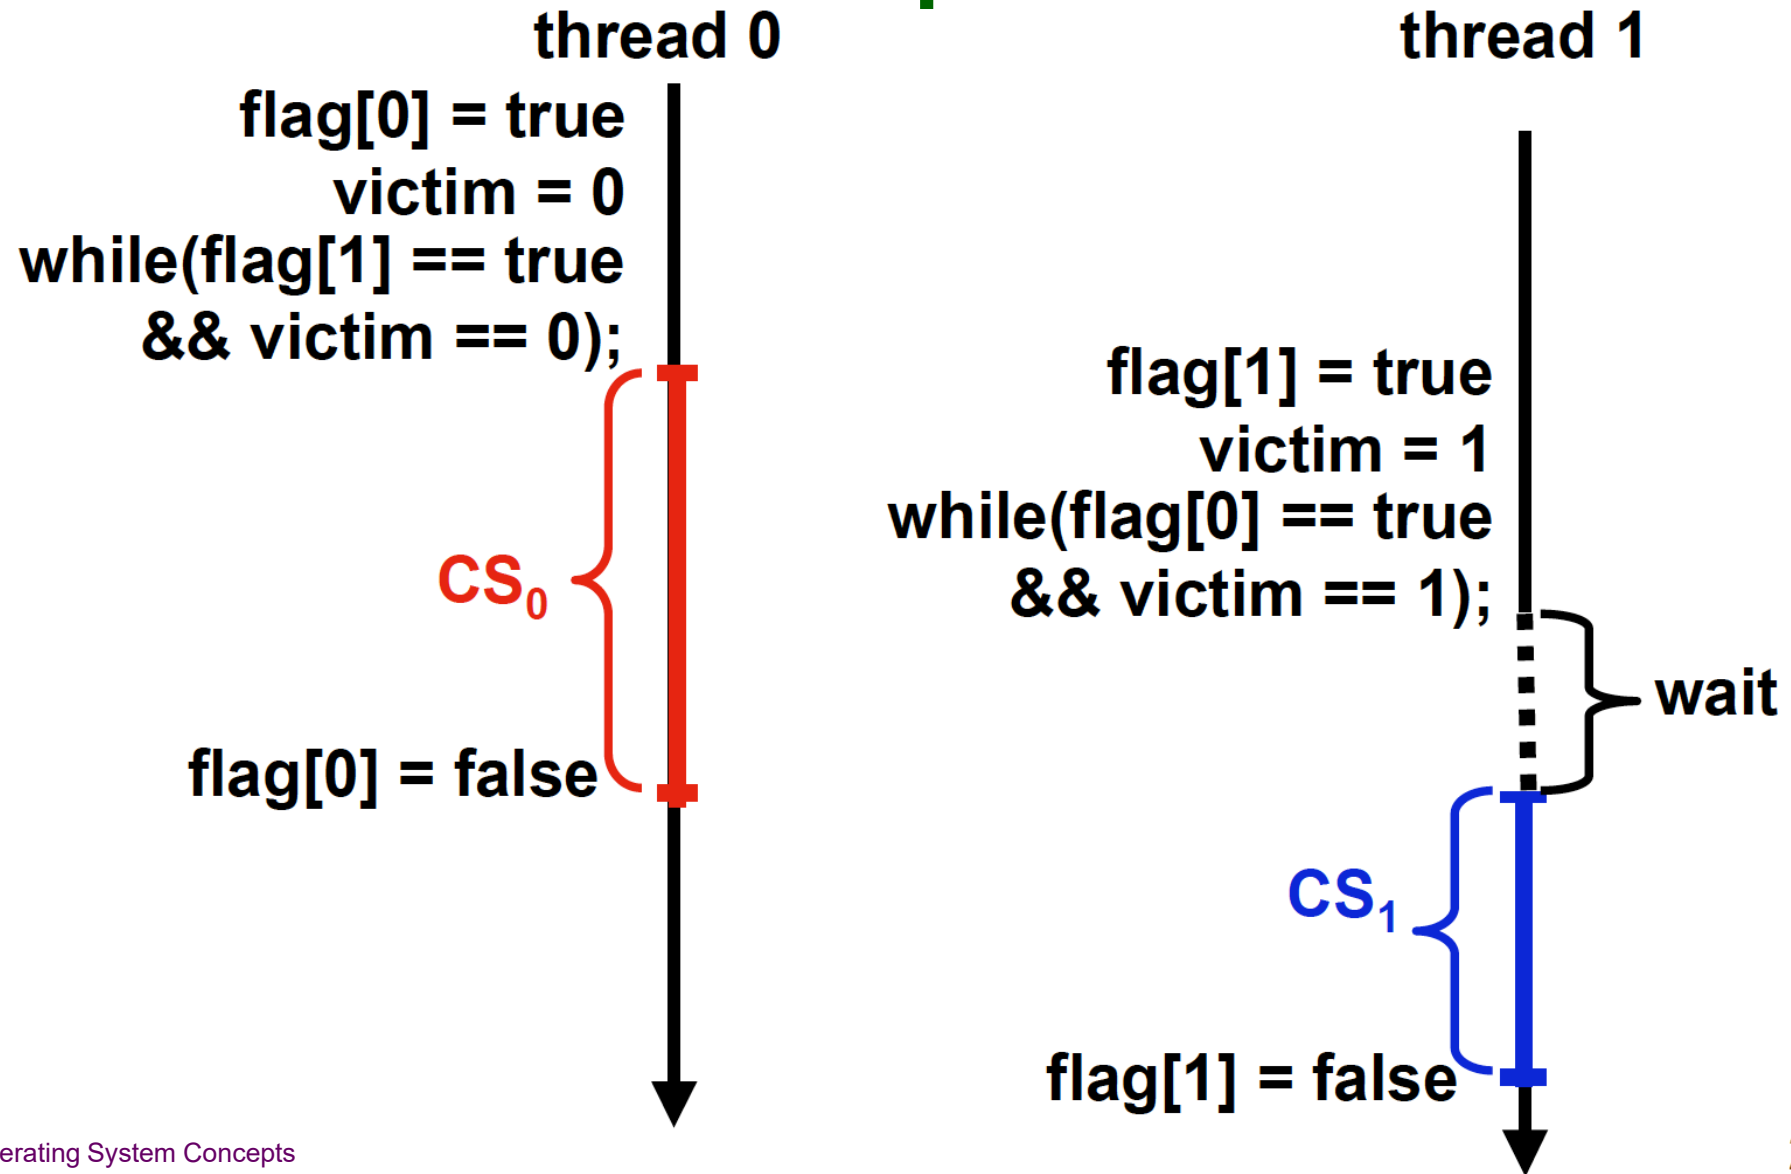
Gary Peterson. Myths about the Mutual Exclusion Problem.  
Information Processing Letters, 12(3):115-116, 1981.

- Meets all the three requirements; solves the critical-section problem for two processes.



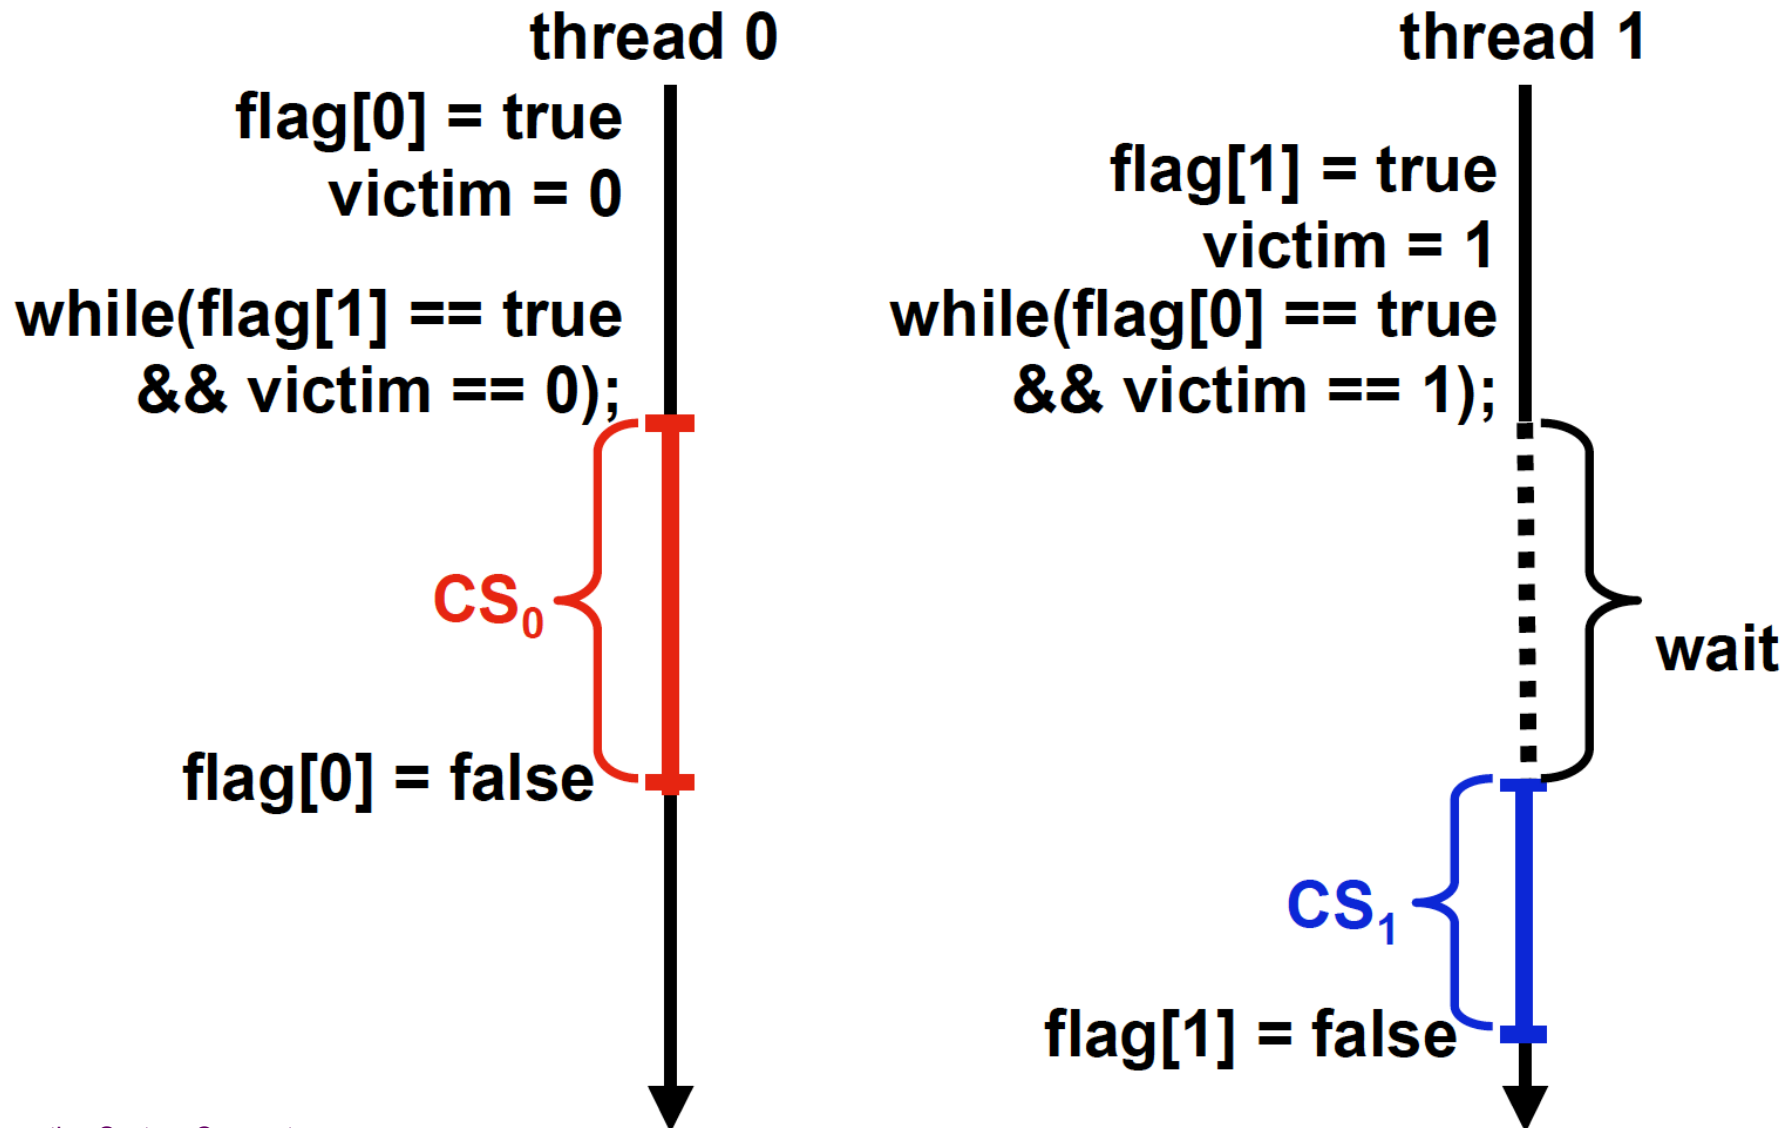



# Peterson's Lock: Serialized Acquires





# Peterson's Lock: Concurrent Acquires






# Proof of Peterson's Algorithm

- The mutual exclusion requirement is assured.
- The progress requirement is assured. The turn variable is only considered when both processes are using, or trying to use, the resource.
- Deadlock is not possible. If both processes are testing the while condition, one of them must have the turn. That process will proceed.
- Finally, bounded waiting is assured. When a process that has exited the CS reenters, it will give away the turn. If the other process is already waiting, it will be the next to proceed.



[https://en.wikipedia.org/wiki/Peterson%27s\\_algorithm](https://en.wikipedia.org/wiki/Peterson%27s_algorithm)




## Is the following code correct?


■ What if we change **victim = i** to **victim = j**?

■ Process  $P_i$

```
do {      flag[i] = true;  // I'm interested
          victim = j;      // I go first
          while (flag[j] and victim == i) ;
              critical section
          flag[i] = false;
              remainder section
      } while (1);
```

■ Does not satisfy mutual exclusion, and not bounded waiting






# Lamport's Bakery Algorithm

## Critical section for $n$ processes

- Before entering its critical section, process receives a number. Holder of the smallest number enters the critical section.
- If processes  $P_i$  and  $P_j$  receive the same number, if  $i < j$ , then  $P_i$  is served first; else  $P_j$  is served first.
- The numbering scheme always generates numbers in increasing order of enumeration; i.e., 1,2,3,3,3,3,4,5...





# Bakery Algorithm

## ■ Notation

◆  $(a,b) < (c,d)$  if  $a < c$  or if  $a = c$  and  $b < d$


◆  $\max(a_0, \dots, a_{n-1})$  is a number,  $k$ , such that  $k \geq a_i$  for  $i = 0, \dots, n-1$


## ■ Shared data

**boolean choosing[n];**

**int number[n];**


Data structures are initialized to **false** and **0** respectively






# Bakery Algorithm


```
do { choosing[i] = true;
    number[i] = max(number[0], number[1], ...,
number [n – 1])+1;
    choosing[i] = false;
    for (j = 0; j < n; j++) {
        while (choosing[j]) ;
        while ((number[j] != 0) && ((number[j],j) <
(number[i],i))) ;
    }
    critical section
    number[i] = 0;
    remainder section
} while (1);
```






# Chapter 6: Process Synchronization


- Background
- The Critical-Section Problem
- Synchronization Hardware
- Semaphores
- Classical Problems of Synchronization
- Monitors
- Synchronization Examples






# Hardware Support

- **There are two types of hardware synchronization supports:**
  - ◆ **Disabling/Enabling interrupts: This is slow and difficult to implement on multiprocessor systems.**
  - ◆ **Special machine instructions:**
    - ✓ **Test and set (TS)**
    - ✓ **Swap**








# Interrupt Disabling


```
do {
```



```
} while (1);
```

- Because interrupts are disabled, no context switch will occur in a critical section.
- Infeasible in a multiprocessor system because all CPUs must be informed.
- Some features that depend on interrupts (e.g., clock) may not work properly.






# Test-and-Set

- Test and modify the content of a word atomically

▪

```
boolean TestAndSet(boolean &target) {  
    boolean rv = target;  
    target = true;  
    return rv;  
}
```





# Mutual Exclusion with Test-and-Set

- Shared data:

**boolean lock = false;**

- Process  $P_i$

**do {**


**while (TestAndSet(lock)) ;**


critical section

**lock = false;**

remainder section

**}**






# Swap

- Atomically swap two variables.

```
void Swap(boolean &a, boolean &b) {  
    boolean temp = a;  
    a = b;  
    b = temp;  
}
```







# Mutual Exclusion with Swap

- Shared data (initialized to **false**):  
**boolean lock;**
- local variable  
**boolean key;**
- Process  $P_i$   
do {  
    **key = true;**  
    **while (key == true)**  
        **Swap(lock, key);**  
    critical section  
    **lock = false;**  
    remainder section

}






# Bounded Waiting Mutual Exclusion with TestAndSet


## Enter Critical Section

```
waiting[i] = true;  
key = true;  
while (waiting[i] && key)  
    key =  
    TestAndSet(lock);  
waiting[i] = false;
```

## Leave Critical Section


```
j = (i+1)%n  
while ((j!=i) && !waiting[j])  
    j = (j+1)%n;  
if (j == i)  
    lock = false;  
else  
    waiting[j] = false;
```






# Chapter 6: Process Synchronization

- Background
- The Critical-Section Problem
- Synchronization Hardware
- **Semaphores**
- Classical Problems of Synchronization
- Monitors
- Synchronization Examples





# Semaphores

- Synchronization tool that does not require busy waiting.
- Semaphore  $S$  – integer variable
- can only be accessed via two indivisible (atomic) operations


*wait (S):*

**while  $S \leq 0$  do *no-op*;**


**$S--$ ;**

*signal (S):*

**$S++$ ;**








# Critical Section of $n$ Processes


- Shared data:

**semaphore mutex;** //initially *mutex* = 1

- Process  $P_i$ :

```
do {  
    wait(mutex);  
    critical section  
    signal(mutex);  
    remainder section  
} while (1);
```







# Semaphore as a General Synchronization Tool

- Execute  $B$  in  $P_j$  only after  $A$  executed in  $P_i$
- Use semaphore *flag* initialized to 0
- Code:

$P_i$	$P_j$
$\vdots$	$\vdots$
$A$	$wait(flag)$
$signal(flag)$	$B$





# Semaphore Implementation




- Define a semaphore as a record


```
typedef struct {  
    int value;  
    struct process *L;  
} semaphore;
```

- Assume two simple operations:

- ◆ **block** suspends the process that invokes it.
- ◆ **wakeup(*P*)** resumes the execution of a blocked process *P*.







# Implementation

- Semaphore operations now defined as

*wait(S):*

**S.value--;**

**if (S.value < 0) {**

add this process to **S.L;**

**block;**

**}**

*signal(S):*


**S.value++;**


**if (S.value <= 0) {**

remove a process **P** from **S.L;**

**wakeup(P);**

**}**







# Deadlock and Starvation

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes.
- Let  $S$  and  $Q$  be two semaphores initialized to 1

$P_0$	$P_1$
$wait(S);$	$wait(Q);$
$wait(Q);$	$wait(S);$
$\vdots$	$\vdots$
$signal(S);$	$signal(Q);$
$signal(Q)$	$signal(S);$


- **Starvation** – indefinite blocking. A process may never be removed from the semaphore queue in which it is suspended.






# Chapter 6: Process Synchronization


- Background
- The Critical-Section Problem
- Synchronization Hardware
- Semaphores
- Classical Problems of Synchronization
- Monitors
- Synchronization Examples






# Classical Problems of Synchronization

- Bounded-Buffer Problem
- Readers and Writers Problem
- Dining-Philosophers Problem








# Bounded-Buffer Problem


- Shared data

**semaphore full, empty, mutex;**

Initially:


**full = 0, empty = n, mutex = 1**






# Bounded-Buffer Problem Producer Process


```
do { ...  
    produce an item in nextp  
    ...  
    wait(empty);  
    wait(mutex);  
    ...  
    add nextp to buffer  
    ...  
    signal(mutex);  
    signal(full);  
} while (1);
```






# Bounded-Buffer Problem Consumer Process

```
do {  
    wait(full)  
    wait(mutex);  
    ...  
    remove an item from buffer to nextc  
    ...  
    signal(mutex);  
    signal(empty);  
    ...  
    consume the item in nextc  
    ...  
} while (1);
```






# Readers-Writers Problem


- Shared data

**semaphore mutex, wrt;**

Initially

**mutex = 1, wrt = 1, readcount = 0**





# Readers-Writers Problem Writer Process


**wait(wrt);**


...

writing is performed

...


**signal(wrt);**







# Readers-Writers Problem Reader Process

```
wait(mutex);  
readcount++;  
if (readcount == 1)  
    wait(wrt);  
signal(mutex);  
  
...  
reading is performed  
  
...  
wait(mutex);  
readcount--;  
if (readcount == 0)  
    signal(wrt);  
signal(mutex);
```






# Dining-Philosophers Problem



■ Shared data

**semaphore chopstick[5];**


Initially all values are 1






# Dining-Philosophers Problem

```
■ Philosopher i:  
    do {  
        wait(chopstick[i])  
        wait(chopstick[(i+1) % 5])  
        ...  
        eat  
        ...  
        signal(chopstick[i]);  
        signal(chopstick[(i+1) % 5]);  
        ...  
        think  
        ...  
    } while (1);
```









# Chapter 6: Process Synchronization

- Background
- The Critical-Section Problem
- Synchronization Hardware
- Semaphores
- Classical Problems of Synchronization
- **Monitors**
- Synchronization Examples







# Monitors

- High-level synchronization construct that allows the safe sharing of an abstract data type among concurrent processes.

**monitor** *monitor-name*


```
{  shared variable declarations
  procedure body  $P_1$  (...) {
    . . .}
  procedure body  $P_2$  (...) {
    . . .}
  procedure body  $P_n$  (...) {
    . . .}
  { initialization code}
}
```








# Monitors: Mutual Exclusion


- *No more than one process* can be executing *within* a monitor. Thus, *mutual exclusion* is guaranteed within a monitor.
- When a process calls a monitor procedure and enters the monitor successfully, it is the *only* process executing in the monitor.
- When a process calls a monitor procedure and the monitor has a process running, the caller will be blocked *outside of the monitor*.









# Schematic View of a Monitor





*processes waiting  
to enter monitor*






# Monitors


- To allow a process to wait within the monitor, a **condition** variable must be declared, as  
**condition x, y;**
- Condition variable can only be used with the operations **wait** and **signal**.
  - ◆ The operation  
**x.wait();**  
means that the process invoking this operation is suspended until another process invokes  
**x.signal();**
  - ◆ The **x.signal** operation resumes exactly one suspended process. If no process is suspended, then the **signal** operation has no effect.






# Condition Signal

- **Consider the released process (from the signaled condition) and the process that signals. There are **two** processes executing in the monitor, and mutual exclusion is violated!**
- There are two common and popular approaches to address this problem:
  - ◆ The released process takes the monitor and the signaling process waits somewhere.
  - ◆ The released process waits somewhere and the signaling process continues to use the monitor.





# Semaphore vs. Condition


Semaphores	Condition Variables
Can be used anywhere, but not in a monitor	Can only be used in monitors
<code>wait()</code> does not always block its caller	<code>wait()</code> <b>always</b> blocks its caller
<code>signal()</code> either releases a process, or increases the semaphore counter	<code>signal()</code> either releases a process, or the signal is <b>lost</b> as if it never occurs
If <code>signal()</code> releases a process, the caller and the released <b>both continue</b>	If <code>signal()</code> releases a process, either the caller or the released continues, but <b>not both</b>


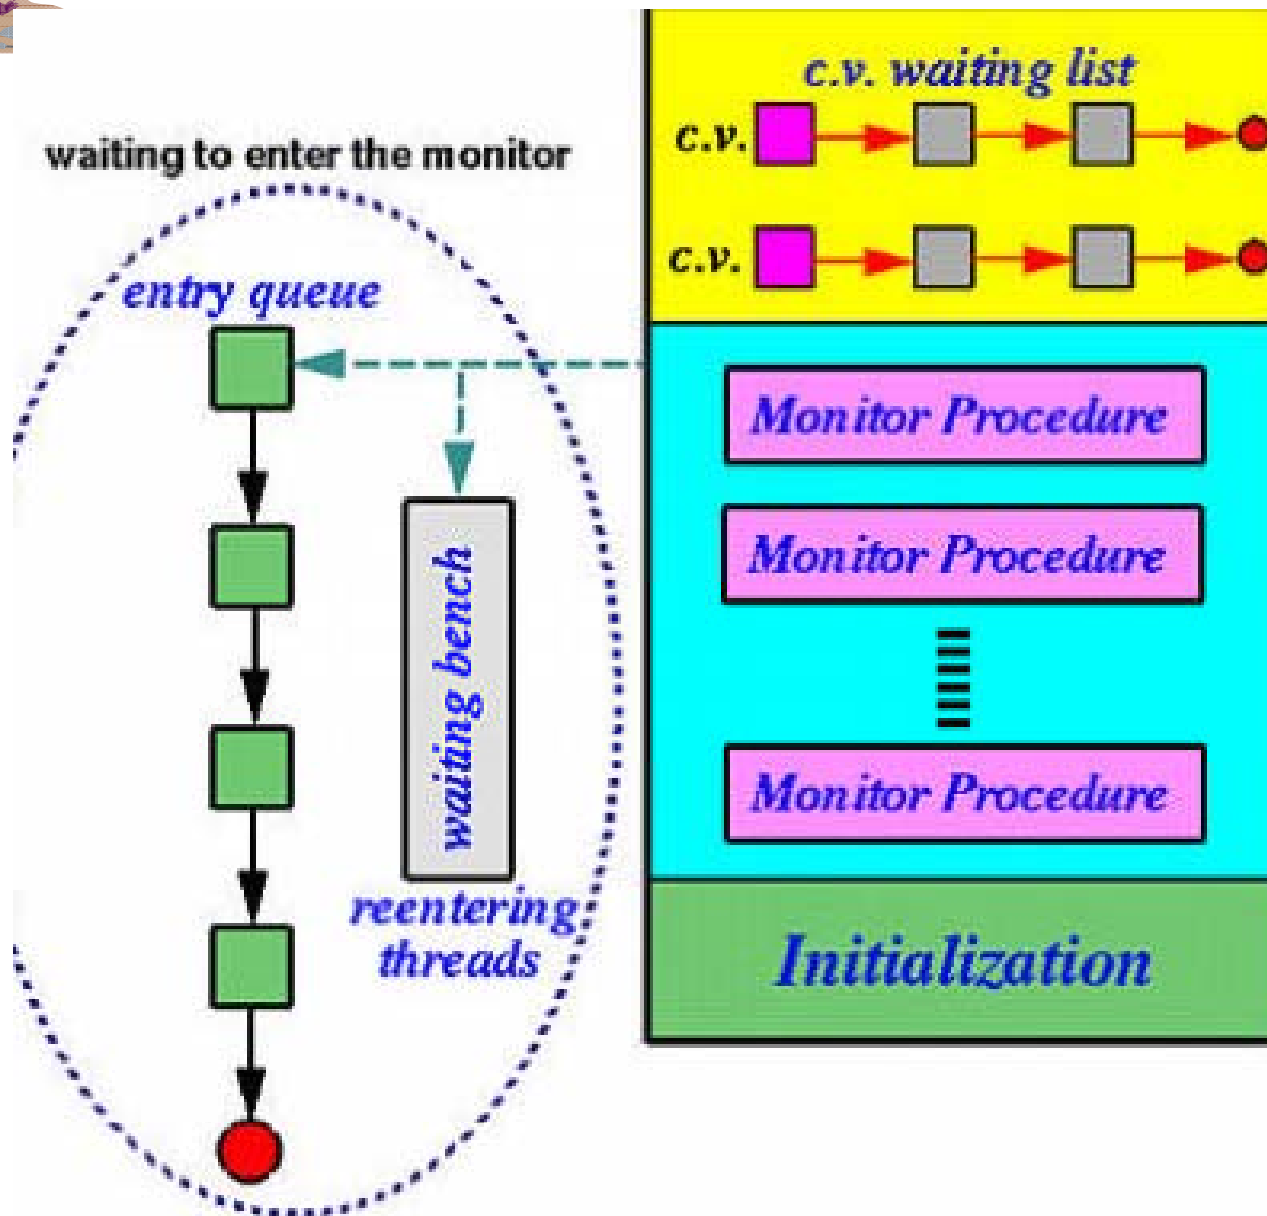
24






# Monitor With Condition Variables







# Dining Philosophers Example

**monitor dp**

**{**

**enum {thinking, hungry, eating} state[5];**

**condition self[5];**

**void pickup(int i) // following slides**


**void putdown(int i) // following slides**


**void test(int i) // following slides**

**void init() { for (int i = 0; i < 5; i++)  
state[i] = thinking;**

**}**

**}**







# Dining Philosophers

```
void pickup(int i) {  
    state[i] = hungry;  
    test(i);  
    if (state[i] != eating)  
        self[i].wait();  
}
```


```
void putdown(int i) {  
    state[i] = thinking;  
    // test left and right neighbors  
    test((i+4) % 5);  
    test((i+1) % 5);  
}
```






# Dining Philosophers

```
void test(int i) {  
    if ( (state[(i + 4) % 5] != eating)  
    &&  
        (state[i] == hungry) &&  
        (state[(i + 1) % 5] != eating)) {  
        state[i] = eating;  
        self[i].signal();  
    }  
}
```





# Monitor Implementation Using Semaphores

- Variables

```
semaphore mutex; // (initially = 1)
semaphore next;  // (initially = 0)
int next-count = 0;
```

- Each external procedure  $F$  will be replaced by  
**wait(mutex);**


...


body of  $F$ ;

...

```
if (next-count > 0)
    signal(next)
else signal(mutex);
```

- Mutual exclusion within a monitor is ensured.







# Monitor Implementation

- For each condition variable **x**, we have:  
**semaphore x-sem; // (initially = 0)**  
**int x-count = 0;**
- The operation **x.wait** can be implemented as:

```
x-count++;  
if (next-count > 0)  
    signal(next);  
else  
    signal(mutex);  
wait(x-sem);  
x-count--;
```







# Monitor Implementation

- The operation **x.signal** can be implemented as:

```
if (x-count > 0) {  
    next-count++;  
    signal(x-sem);  
    wait(next);  
    next-count--;  
}
```









# Monitor Implementation


- **Conditional-wait** construct: **x.wait(c);**
  - ◆ **c** – integer expression evaluated when the **wait** operation is executed.
  - ◆ value of **c** (a *priority number*) stored with the name of the process that is suspended.
  - ◆ when **x.signal** is executed, process with smallest associated priority number is resumed next.






## Monitor Implementation (Cont.)


- Check two conditions to establish correctness of system:
  - ◆ User processes must always make their calls on the monitor in a correct sequence.
  - ◆ Must ensure that an uncooperative process does not ignore the mutual-exclusion gateway provided by the monitor, and try to access the shared resource directly, without using the access protocols.






# Chapter 6: Process Synchronization

- Background
- The Critical-Section Problem
- Synchronization Hardware
- Semaphores
- Classical Problems of Synchronization
- Monitors
- Synchronization Examples






## Solaris 2 Synchronization

- Implements a variety of locks to support multitasking, multithreading (including real-time threads), and multiprocessing.
- Uses *adaptive mutexes* for efficiency when protecting data from short code segments.
- Uses *condition variables* , *semaphore*, and *readers-writers locks* when longer sections of code need access to data.
- Uses *turnstiles* to order the list of threads waiting to acquire either an adaptive mutex or reader-writer lock.





# Windows XP Synchronization

- Uses interrupt masks to protect access to global resources on uniprocessor systems.
- Uses *spinlocks* on multiprocessor systems.
- Also provides *dispatcher objects* which may act as mutexes and semaphores.
- Dispatcher objects may also provide *events*. An event acts much like a condition variable.

