

Chapter 9: Virtual Memory


肖 卿 俊

办公室：计算机楼532室

电邮： csqjxiao@seu.edu.cn

主页： <http://cse.seu.edu.cn/PersonalPage/csqjxiao>


电话： 025-52091023



Objectives


- To describe the **benefits** of a virtual memory system
- To explain the concepts of **demand paging**, **page-replacement** algorithms, and allocation of page frames
- To discuss the principle of **working-set model**
- To explain the IPC model based on memory sharing; To examine the differences between **shared memory** and **memory-mapped files**
- To explore **how kernel memory is managed**






Chapter 9: Virtual Memory


- Background
- Demand Paging
- Copy-on-Write
- Page Replacement
- Allocation of Frames
- Thrashing
- Memory-Mapped Files
- Allocating Kernel Memory
- Other Considerations
- Operating-System Examples






Background


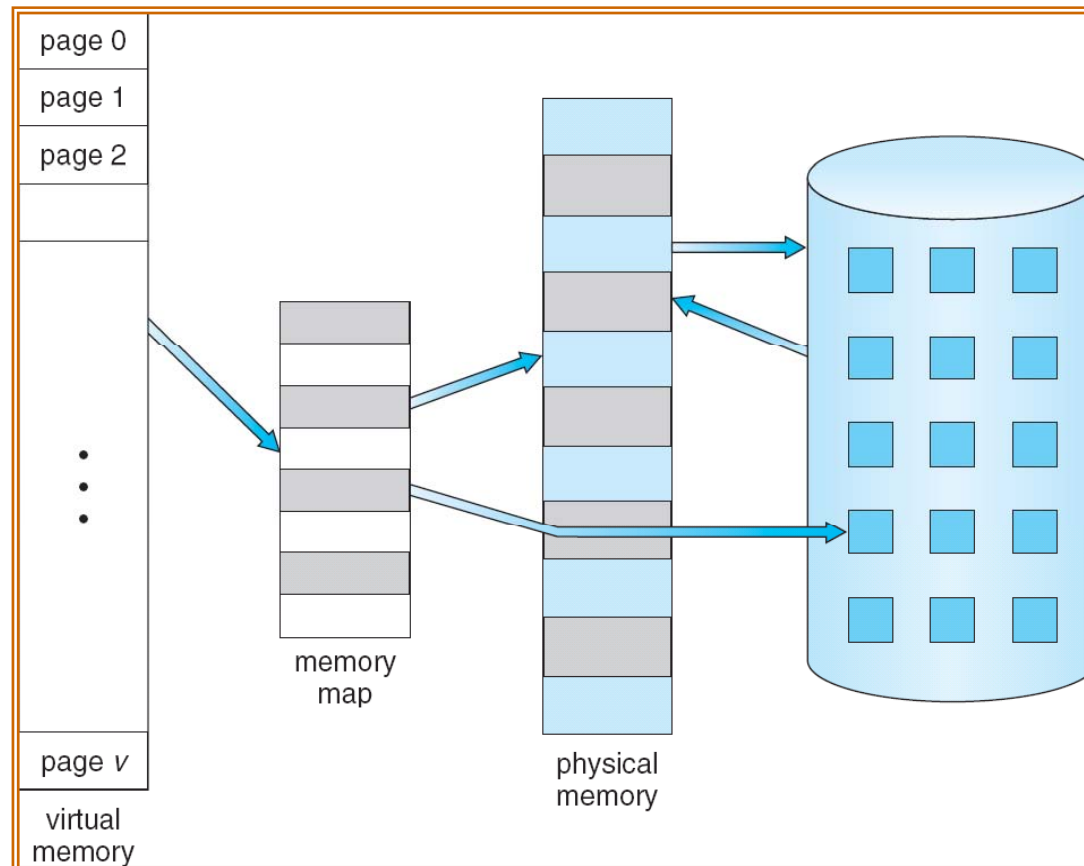
- We previously talked about an entire process swapping into or out of main memory
- **Idea of virtual memory** – separation of user logical memory from physical memory.
 - ◆ Only part of the program needs to be kept in main memory for execution.
 - ◆ Logical address space can therefore be much larger than physical address space.
 - ◆ More programs can be run at the same time
 - ◆ Less I/O is needed than loading or swapping








Two Kinds of Implementation for Virtual Memory


- Virtual memory can be implemented via:
 - Demand paging (按需调页)
 - Demand segmentation (按需调段)




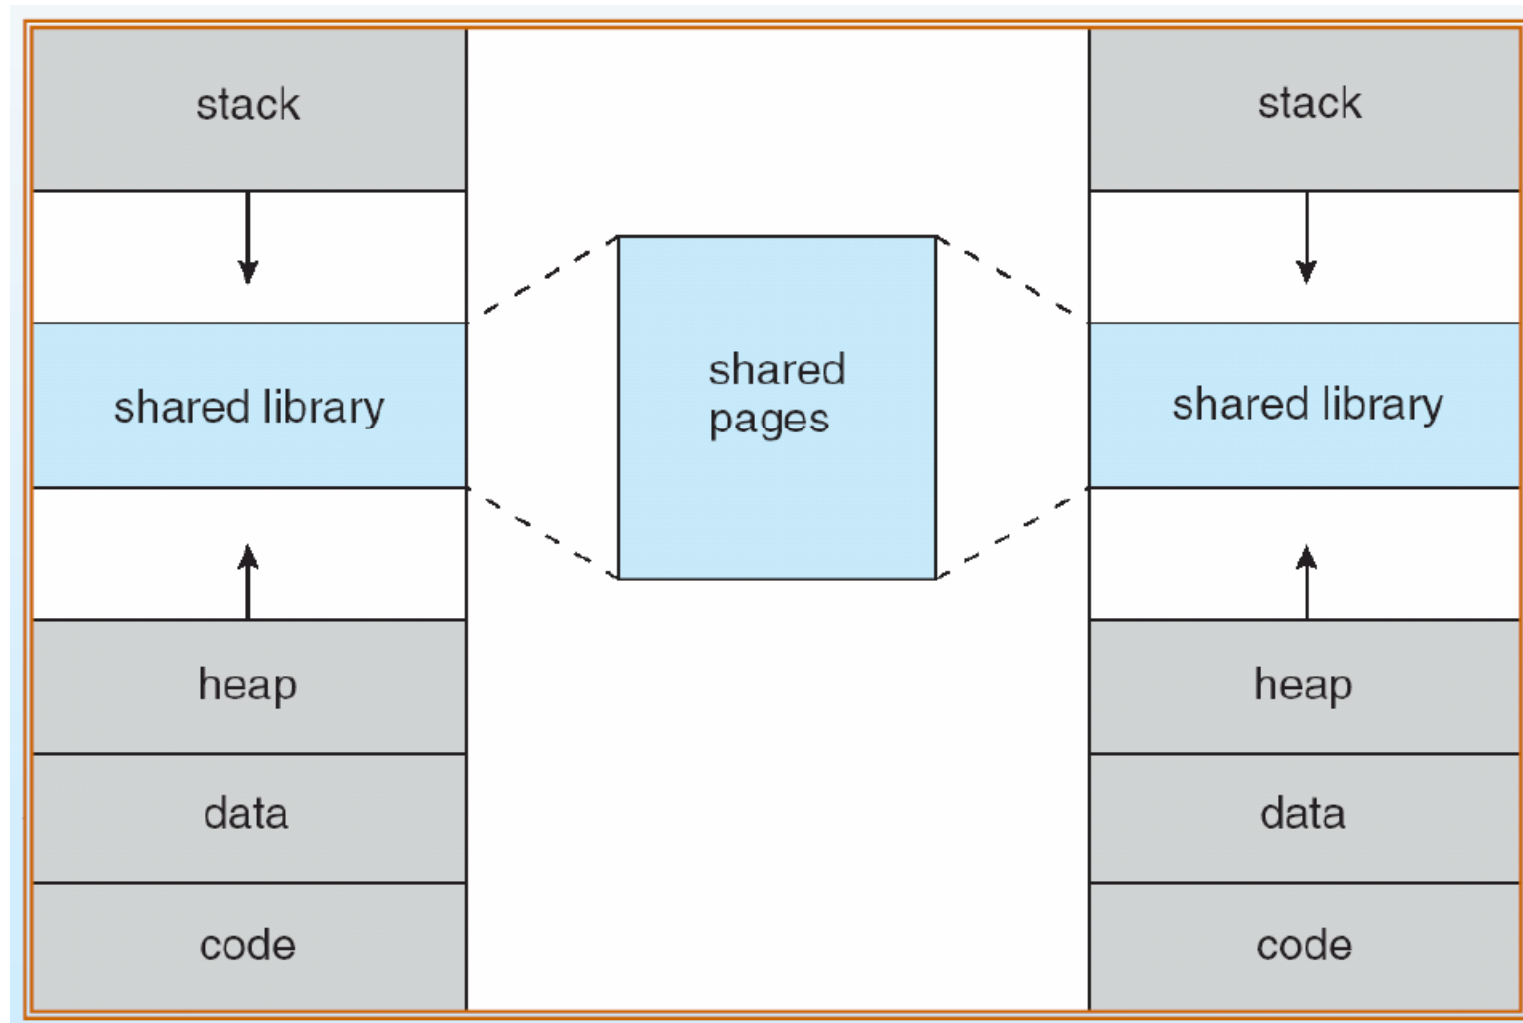



Virtual-Address Space






Shared Library Using Virtual Memory






Chapter 9: Virtual Memory

- Background
- Demand Paging
- Copy-on-Write
- Page Replacement
- Allocation of Frames
- Thrashing
- Memory-Mapped Files
- Allocating Kernel Memory
- Other Considerations
- Operating-System Examples






Demand Paging


- Bring a page into memory only when it is needed.
 - ◆ Less I/O needed
 - ◆ Less memory needed
 - ◆ Faster response
 - ◆ More users
- Page is needed \Rightarrow reference to it
 - ◆ invalid reference \Rightarrow abort
 - ◆ not-in-memory \Rightarrow bring to memory
- **Pure demand paging**– never bring a page into memory unless page will be needed






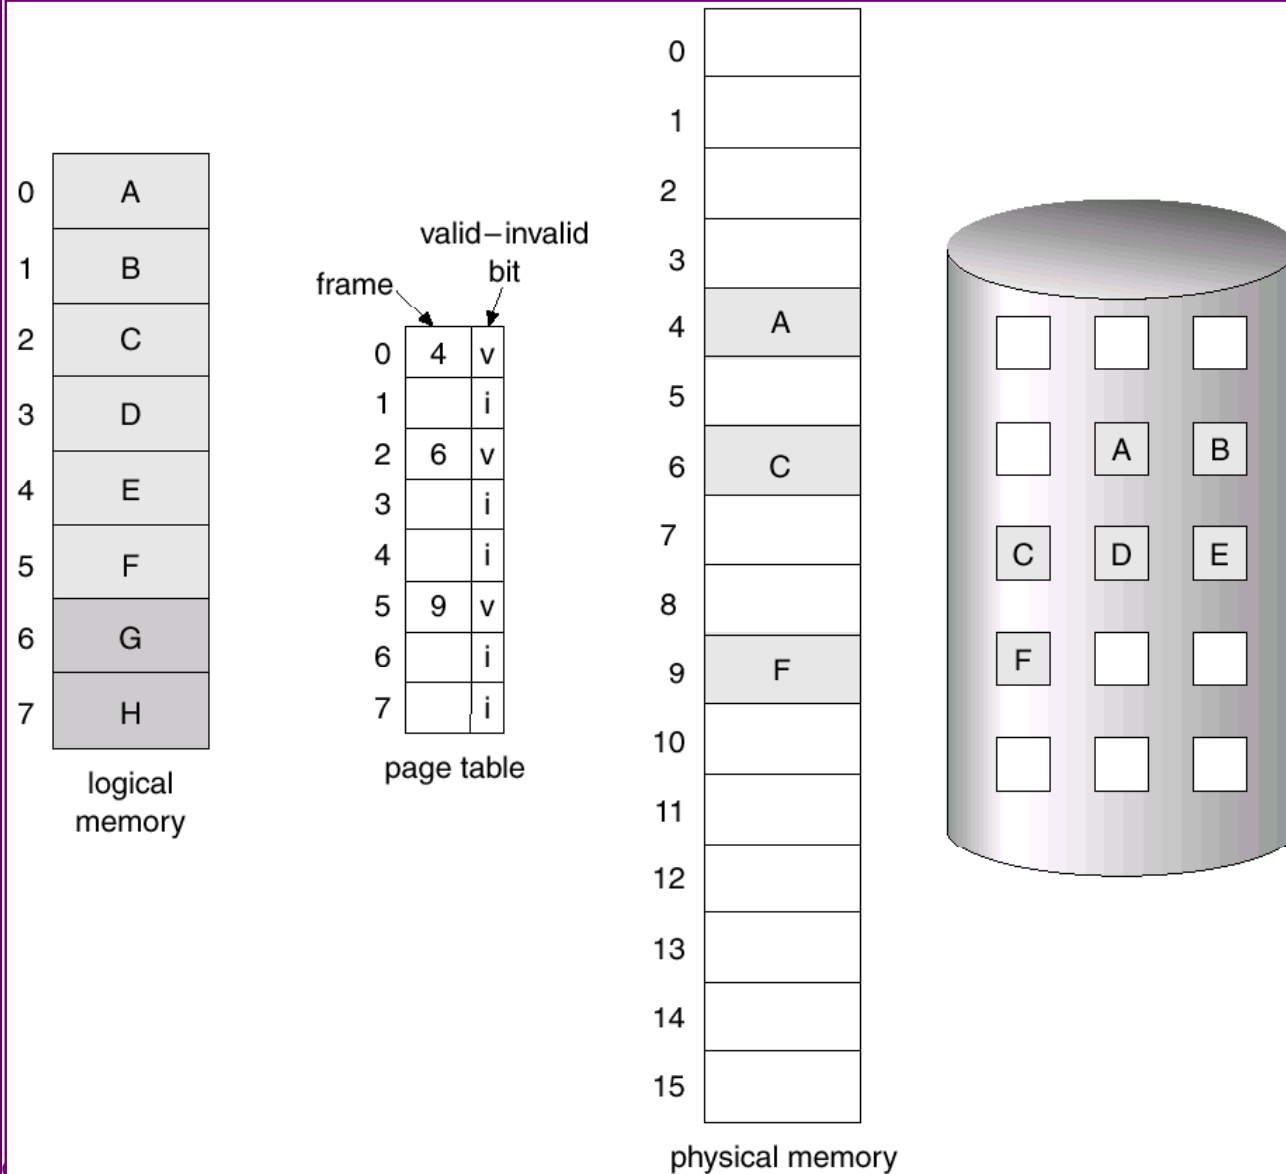
Valid-Invalid Bit


- With each page table entry, a valid-invalid bit is associated
 - ◆ (1 \Rightarrow in-memory, 0 \Rightarrow not-in-memory)
- Initially, valid-invalid bit is set to 0 on all entries.
- During address translation, if valid-invalid bit in page table entry is 0 \Rightarrow page fault.






Page Table When Some Pages Are Not in Main Memory







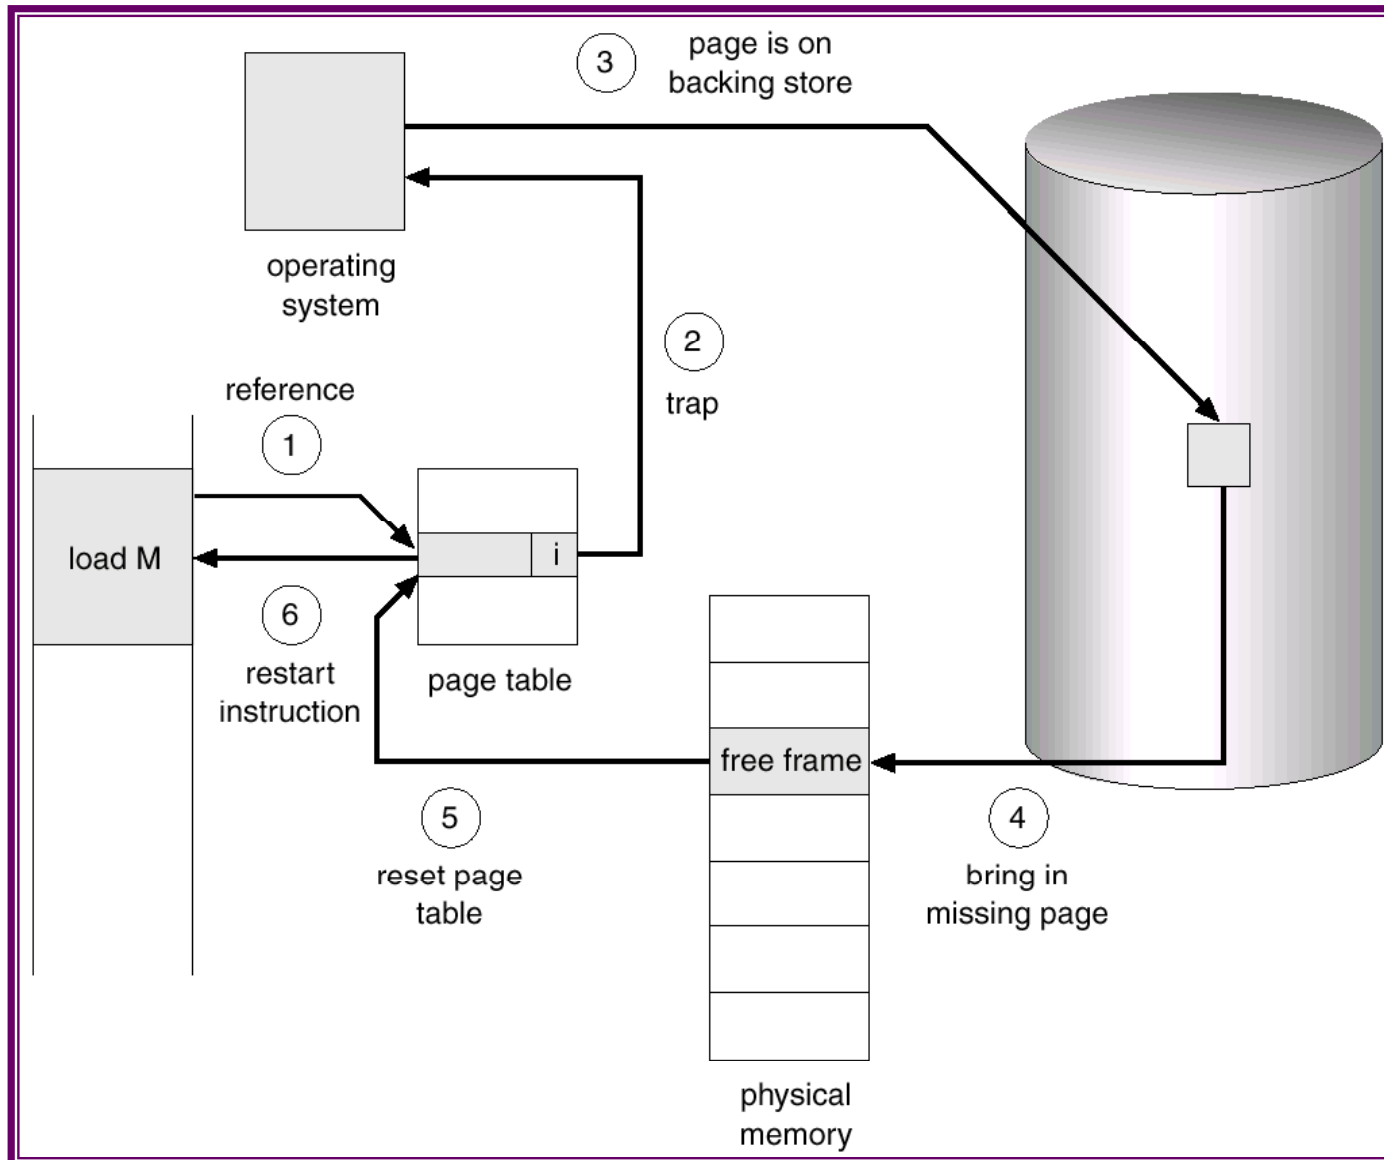
Steps in Handling a Page Fault


- If there is ever a reference to a page, first reference will trap to OS \Rightarrow page fault
- OS looks at another table to decide:
 - ◆ Invalid reference \Rightarrow abort.
 - ◆ Just not in memory.
- Get empty frame.
- Swap page into frame.
- Reset tables, validation bit = 1.
- Restart instruction







Steps in Handling a Page Fault






More Details about Restarting an Instruction

- The restart will require fetching instruction again, decoding it again, fetching the two operands again, and applying it again
- Difficulty arises when an instruction may modify multiple virtual pages
 - ◆ For example, block move operation
 - ◆ Auto increment/decrement location
 - ◆ Restart the whole operation?
 - ✓ What if source and destination overlap?
 - ✓ The source may have been modified





Performance of Demand Paging


■ Page Fault Rate $0 \leq p \leq 1.0$


◆ if $p = 0$, no page faults

◆ if $p = 1$, every reference is a fault

■ Effective Access Time (EAT)


$$\begin{aligned} \text{EAT} = & (1 - p) \times \text{memory access} \\ & + p \times (\text{page fault overhead} \\ & \quad [+ \text{swap page out}] \\ & \quad + \text{swap page in} \\ & \quad + \text{restart overhead}) \end{aligned}$$






Demand Paging Example


- Memory access time = 1 microsecond
- Swap Page Time = 10 millisec = 10000 microseconds
- Assume 50% of the time the page that is being replaced has been modified and therefore needs to be swapped out.
- Ignore the cost of restarting an instruction.
- $$\begin{aligned} \text{EAT} &= (1 - p) \times 1 + p \times (10000 \times 50\% + 20000 \times 50\%) \\ &= (1 - p) \times 1 + p \times (15000) \\ &= 1 + 14999 \times p \quad (\text{in microsecond}) \end{aligned}$$






Chapter 9: Virtual Memory


- Background
- Demand Paging
- Copy-on-Write
- Page Replacement
- Allocation of Frames
- Thrashing
- Memory-Mapped Files
- Allocating Kernel Memory
- Other Considerations
- Operating-System Examples






Copy-on-Write


- **Copy-on-Write** (COW) allows both parent and child processes to initially ***share*** the same pages in memory
 - ◆ If either process modifies a shared page, only then is the page copied
- COW allows more efficient process creation as only modified pages are copied






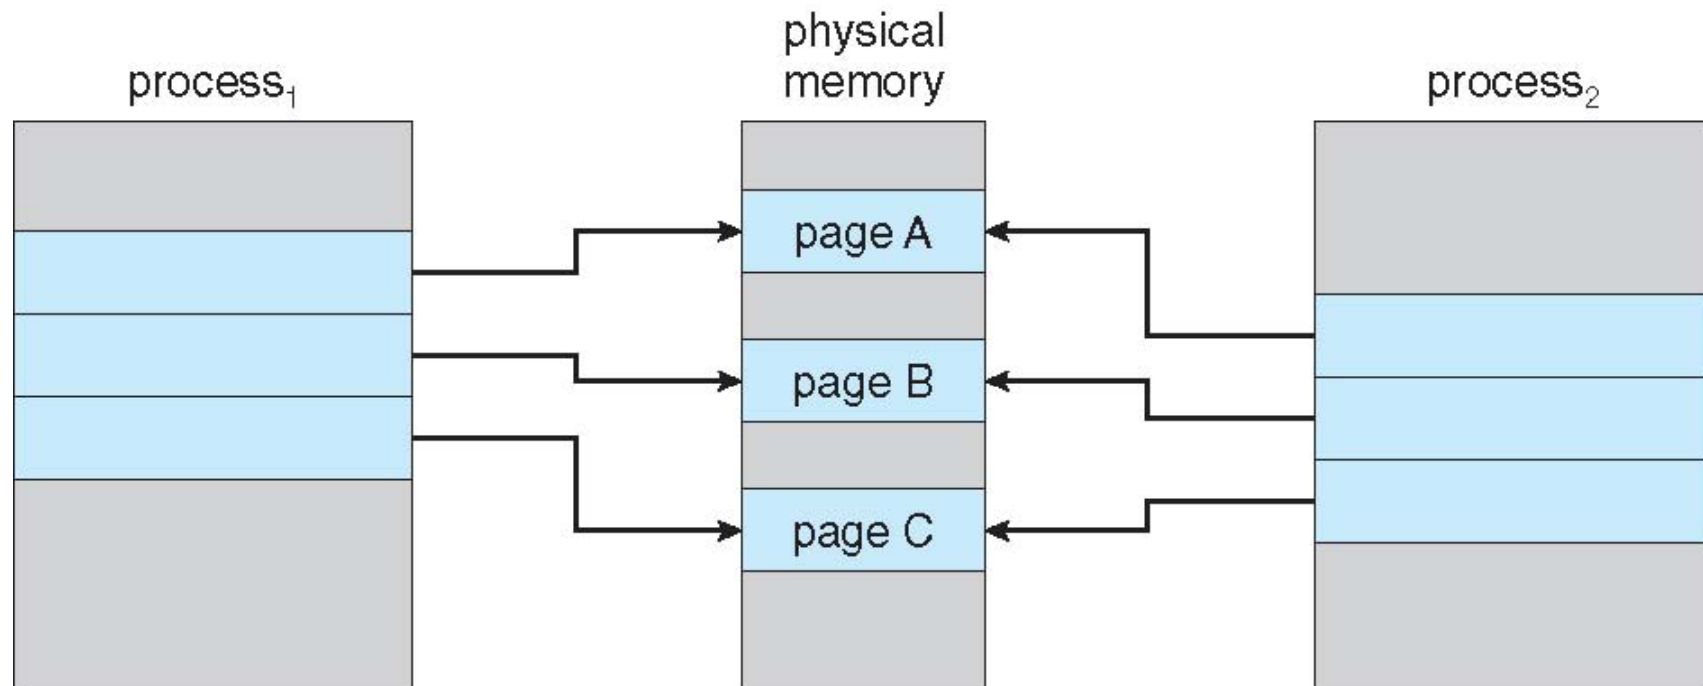
Copy-on-Write (Cont.)


- In general, free pages are allocated from a **pool** of **zero-fill-on-demand** pages
 - ◆ Pool should always have free frames for fast demand page execution
 - ◆ Why zero-out a page before allocating it?
- `vfork()` variation on `fork()` system call has parent suspend and child using copy-on-write address space of parent
 - ◆ Designed to have child call `exec()`
 - ◆ Very efficient




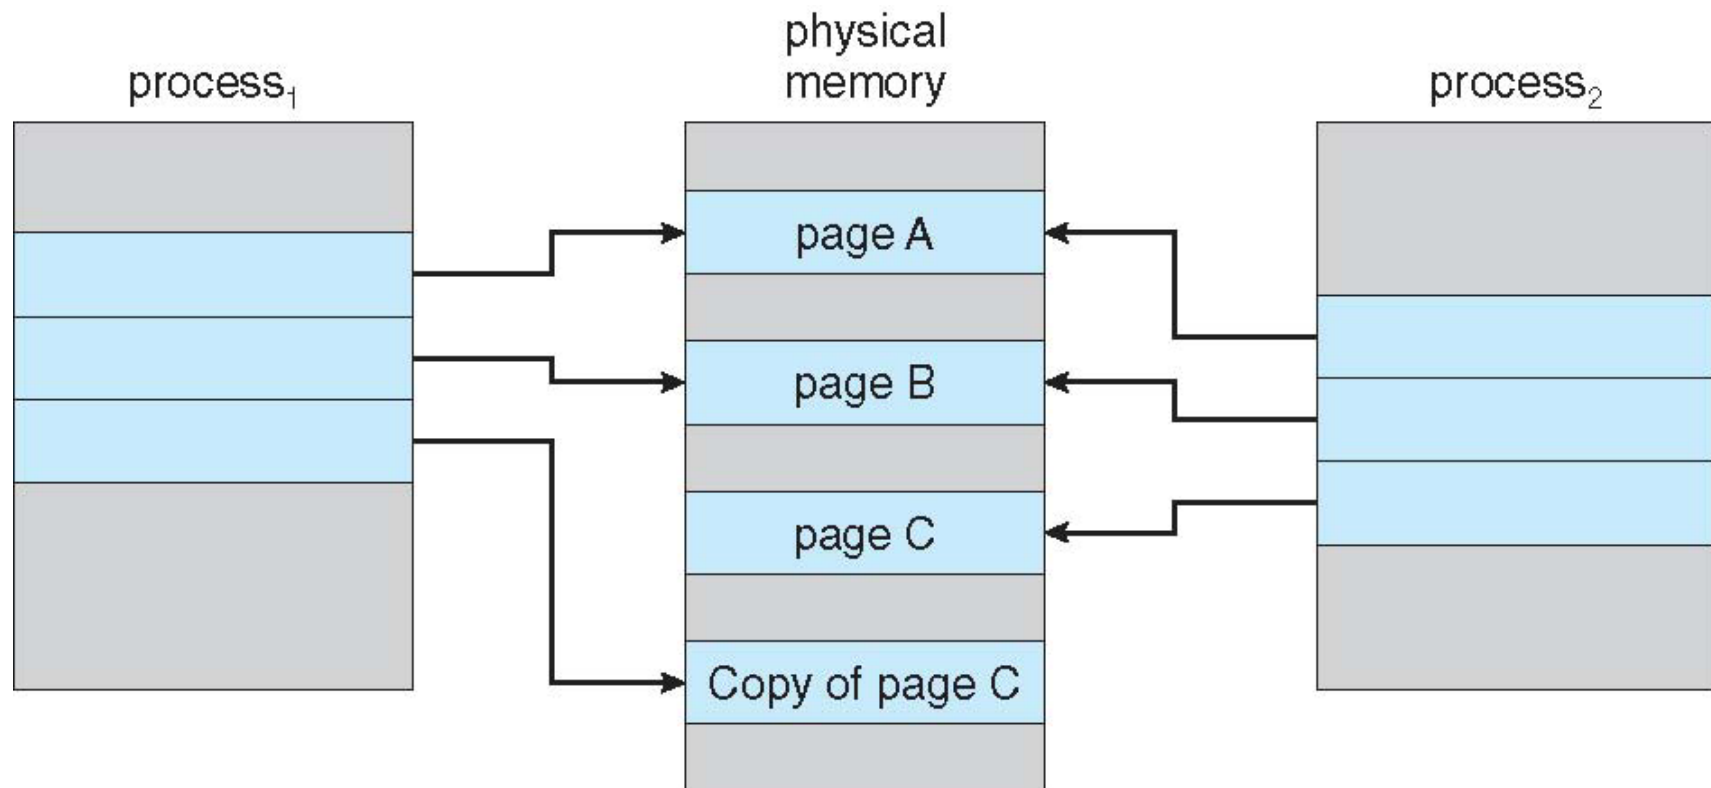



Before Process 1 Modifies Page C






After Process 1 Modifies Page C






Chapter 9: Virtual Memory


- Background
- Demand Paging
- Copy-on-Write
- Page Replacement
- Allocation of Frames
- Thrashing
- Memory-Mapped Files
- Allocating Kernel Memory
- Other Considerations
- Operating-System Examples






What Happens if There is no Free Frame?


- Used up by process pages
- Also in demand by the kernel, I/O buffers, ...
- How much to allocate to each?
- Same page may be brought into memory several times
- Page replacement – find some page in memory, but not really in use, swap it out
 - ◆ Algorithm – terminate? swap out? replace the page?
 - ◆ Performance – want an algorithm which will result in minimum number of page faults







Page Replacement

- Prevent over-allocation of memory by modifying page-fault service routine to include page replacement.
- Use *modify (dirty) bit* to reduce overhead of page transfers – only modified pages are written to disk.
- Page replacement completes separation between logical memory and physical memory – large virtual memory can be provided on a smaller physical memory.




Need For Page Replacement







Basic Page Replacement

1. Find the location of the desired page on disk.
2. Find a free frame:
 - If there is a free frame, use it.
 - If there is no free frame, use a page replacement algorithm to select a *victim* frame and swap it out
3. Read the desired page into the (newly) free frame. Update the page and frame tables.
4. Restart the instruction.




Page Replacement







Page Replacement Algorithms

- Want lowest page-fault rate.
- Evaluate algorithm by running it on a particular string of memory references (reference string) and computing the number of page faults on that string.
- In all our examples, the reference string is
1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5.





Graph of Page Faults Versus The Number of Frames






First-In-First-Out (FIFO) Algorithm

- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- 3 frames (3 pages can be in memory at a time per process)

1	1	4	5
2	2	1	3
3	3	2	4

9 page faults





First-In-First-Out (FIFO) Algorithm


■ 4 frames


1	1	5	4
2	2	1	5
3	3	2	
4	4	3	

10 page faults

■ FIFO Replacement – Belady’s Anomaly

◆ more frames \Rightarrow less page faults







FIFO Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

3 page frames





FIFO Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1


7	7	7	2
	0	0	0
		1	1

2	2	4	4	4	0
3	3	3	2	2	2
1	0	0	0	3	3

0	0
1	1
3	2



7	7	7
1	0	0
2	2	1


3 page frames





FIFO Illustrating Belady's Anomaly






Optimal Algorithm


- Replace page that will not be used for longest period of time.
- 4 frames example


1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5



6 page faults

- How do you know this?
- Used for measuring how well your algorithm performs.






Optimal Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		2		2		2						7
	0	0	0		0		0		0		0						0
		1	1		3		3		3		1						1

page frames





Least Recently Used (LRU) Algorithm


- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

1	5
2	
3	5 4
4	3

- Counter implementation

- ◆ Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter.
- ◆ When a page needs to be changed, look at the counters to determine which are to change.






LRU Page Replacement


reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		4	4	4	0			1		1		1		
	0	0	0		0		0	0	3	3			3		0		0		
		1	1		3		3	2	2	2			2		2		7		


page frames





LRU Algorithm (Cont.)

- Stack implementation – keep a stack of page numbers in a double link form:
 - ◆ Page referenced:
 - ✓ move it to the top
 - ✓ requires 6 pointers to be changed
 - ◆ No search for replacement





Use Of A Stack to Record The Most Recent Page References

reference string

4 7 0 7 1 0 1 2 1 2 7 1 2

2
1
0
7
4

stack before a


7
2
1
0
4

stack after b

↑
a


↑
b






Problems of Previous LRU Implementations

- Two implementations of LRU
 - ◆ Clock: Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter.
 - ◆ Stack: Whenever a page is referenced, it is removed from the stack and put on the top.
- The updating of the clock fields or stack must be done for every memory reference
- Would slow every memory access by a factor of at least ten







LRU Approximation Algorithms

- Reference bit (Hardware maintained)
 - ◆ Each page is associated with a bit in the page table
 - ◆ Initially 0; When page is referenced, set the bit to 1.
 - ◆ Replace the one which is 0 (if one exists)

- However, we do not know the order of use.


- This information is the basis for many page-replacement algorithms that approximate LRU replacement






LRU Approximation Algorithms


- Rational: Gain additional ordering information by recording the reference bits at regular intervals
- Additional-Reference-Bits Algorithm
 - ◆ Keep an 8-bit bytes for each page in main memory
 - ◆ At regular intervals, shifts the bits right 1 bit, shift the reference bit into the high-order bit
 - ◆ Interpret these 8-bit bytes as unsigned intergers, the page with lowest number is the LRU page






LRU Approximation Algorithms


- Second-Chance Algorithm (FIFO+reference bit)
 - ◆ When a page has been selected for replacement, we inspect its reference bit.
 - ◆ If the value is 0, we proceed to replace this page;
 - ◆ but if the reference bit is set to 1, we give the page a second chance and move on to select the next FIFO page.
 - ◆ When a page gets a second chance, its reference bit is cleared, and its arrival time is reset to the current time.



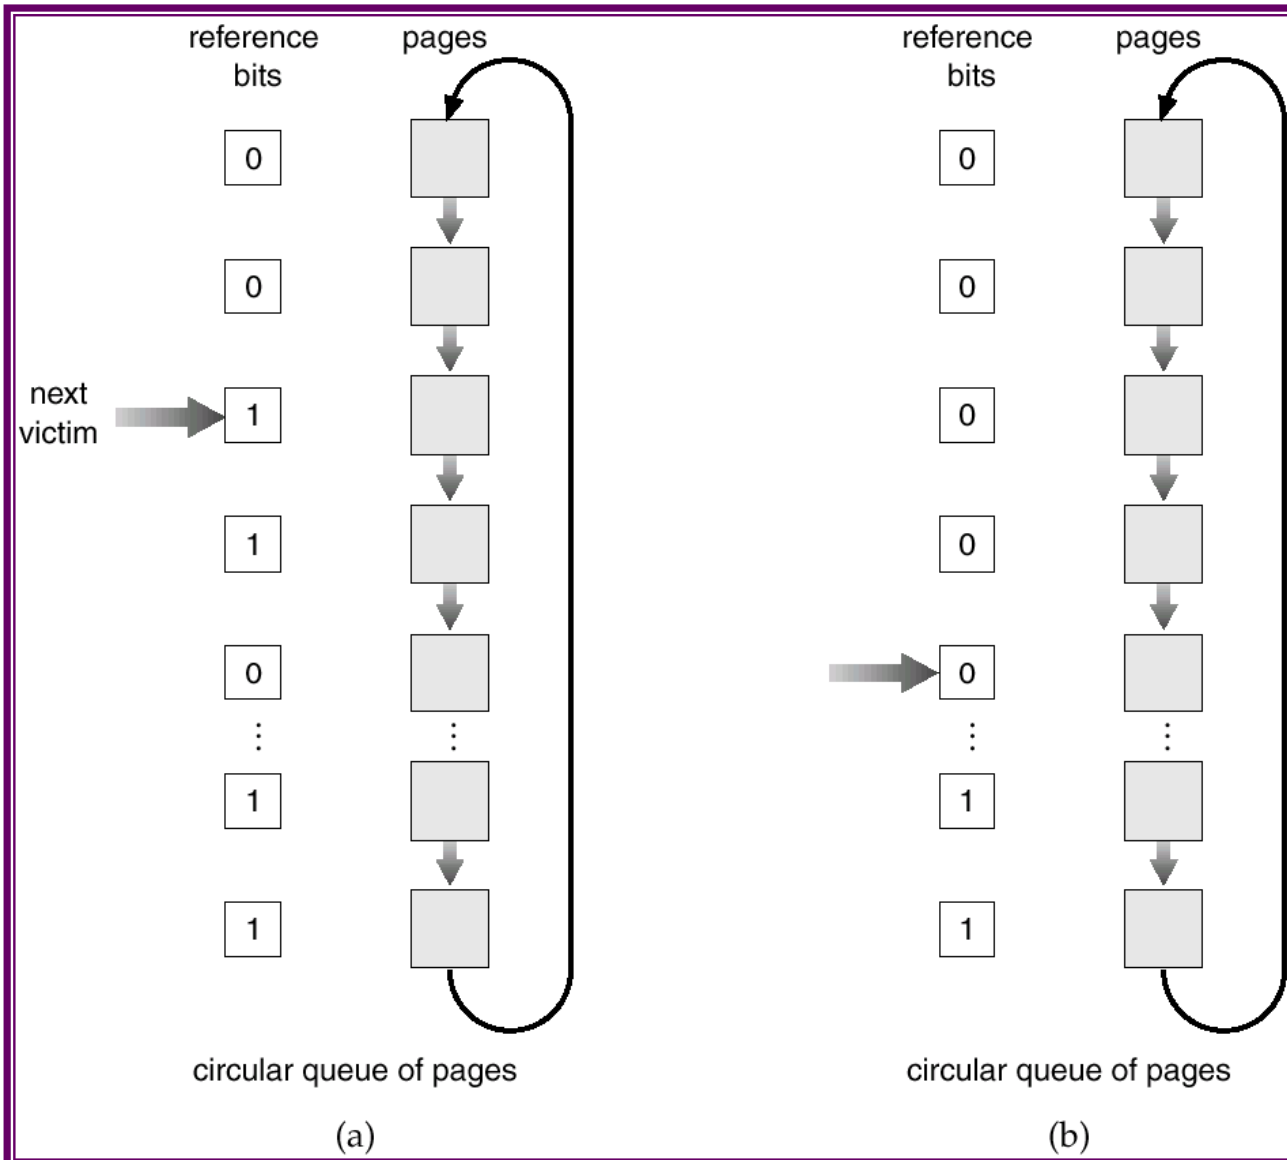



LRU Approximation Algorithms

- Second-Chance Algorithm (clock+reference bit)
 - ◆ Given a circular queue, called clock
 - ◆ If page to be replaced (in clock order) has reference bit = 1. then:
 - ✓ set reference bit 0.
 - ✓ leave page in memory.
 - ✓ replace next page (in clock order), subject to same rules.




Second-Chance (clock) Page-Replacement Algorithm






Counting Algorithms


- Keep a counter of the number of references that have been made to each page.
- **Least Frequently Used (LFU) Algorithm:** replaces page with smallest count.
- **Most Frequently Used (MFU) Algorithm:** based on the argument that the page with the smallest count was probably just brought in and has yet to be used.






Chapter 9: Virtual Memory


- Background
- Demand Paging
- Copy-on-Write
- Page Replacement
- Allocation of Frames
- Thrashing
- Memory-Mapped Files
- Allocating Kernel Memory
- Other Considerations
- Operating-System Examples






Allocation of Frames

- Each process needs **minimum** number of pages.
- Example: IBM 370 – 6 pages to handle SS MOVE instruction:
 - ◆ instruction is 6 bytes, might span 2 pages.
 - ◆ 2 pages to handle **from**.
 - ◆ 2 pages to handle **to**.
- Two major allocation schemes.
 - ◆ fixed allocation
 - ◆ priority allocation





Fixed Allocation


- Equal allocation – e.g., if 100 frames and 5 processes, give each 20 pages.
- Proportional allocation – Allocate according to the size of process.


s_i = size of process p_i

$$S = \sum s_i$$

m = total number of frames


$$a_i = \text{allocation for } p_i = \frac{s_i}{S} \times m$$






Priority Allocation


- Use a proportional allocation scheme using priorities rather than size.
- If process P_i generates a page fault,
 - ◆ select for replacement one of its frames.
 - ◆ select for replacement a frame from a process with lower priority number.






Global vs. Local Allocation


- **Global** replacement – process selects a replacement frame from the set of all frames; one process can take a frame from another.
- **Local** replacement – each process selects from only its own set of allocated frames.






Chapter 9: Virtual Memory

- Background
- Demand Paging
- Copy-on-Write
- Page Replacement
- Allocation of Frames
- **Thrashing**
- Memory-Mapped Files
- Allocating Kernel Memory
- Other Considerations
- Operating-System Examples







Thrashing


- If a process does not have “enough” frames, the page-fault rate is very high. This leads to:
 - ◆ low CPU utilization.
 - ◆ operating system thinks that it needs to increase the degree of multiprogramming.
 - ◆ another process added to the system.

- **Thrashing** \equiv a process is busy swapping pages in and out.





Thrashing




■ Why does paging work?


Locality model

- ◆ Process migrates from one locality to another.
- ◆ Localities may overlap.



■ Why does thrashing occur?


Σ size of locality > total memory size






Locality In A Memory-Reference Pattern






Working-Set Model


- $\Delta \equiv$ working-set window \equiv a fixed number of page references
Example: 10,000 instruction
- WSS_i (working set of Process P_i) = total number of pages referenced in the most recent Δ (varies in time)
 - ◆ if Δ too small will not encompass entire locality.
 - ◆ if Δ too large will encompass several localities.
 - ◆ if $\Delta = \infty \Rightarrow$ will encompass entire program.






Working-Set Model (Cont.)

- $D = \sum WSS_i \equiv$ total demand frames
- if $D > m \Rightarrow$ Thrashing
- Policy if $D > m$, then suspend one of the processes.







Working-Set Model

page reference table


... 2 6 1 5 7 7 7 5 1 6 2 3 4 1 2 3 4 4 4 3 4 3 4 4 4 1 3 2 3 4 4 4 3 4 4 4 ...




$$WS(t_1) = \{1, 2, 5, 6, 7\}$$




$$WS(t_2) = \{3, 4\}$$







Keeping Track of the Working Set

- Approximate with interval timer + a reference bit
- Example: $\Delta = 10,000$
 - ◆ Timer interrupts after every 5000 time units.
 - ◆ Keep in memory 2 bits for each page.
 - ◆ Whenever a timer interrupts copy and sets the values of all reference bits to 0.
 - ◆ If one of the bits in memory = 1 \Rightarrow page in working set.
- Why is this not completely accurate?
- Improvement = 10 bits and interrupt every 1000 time units.







Page-Fault Frequency Scheme




- Establish “acceptable” page-fault rate.
 - ◆ If actual rate too low, process loses frame.
 - ◆ If actual rate too high, process gains frame.






Chapter 9: Virtual Memory

- Background
- Demand Paging
- Copy-on-Write
- Page Replacement
- Allocation of Frames
- Thrashing
- **Memory-Mapped Files**
- Allocating Kernel Memory
- Other Considerations
- Operating-System Examples








Memory-Mapped Files

- Memory-mapped file I/O allows file I/O to be treated as routine memory access by *mapping* a disk block to a page in memory.
- A file is initially read using demand paging. A page-sized portion of the file is read from the file system into a physical page. Subsequent reads/writes to/from the file are treated as ordinary memory accesses.
- Simplifies file access by treating file I/O through memory rather than **read()** **write()** system calls.
- Also allows several processes to map the same file allowing the pages in memory to be shared.








Memory-Mapped Files






Memory-Mapped Shared Memory in Windows






Chapter 9: Virtual Memory


- Background
- Demand Paging
- Copy-on-Write
- Page Replacement
- Allocation of Frames
- Thrashing
- Memory-Mapped Files
- Allocating Kernel Memory
- Other Considerations
- Operating-System Examples






Allocating Kernel Memory

- Treated differently from user memory
- Often allocated from a free-memory pool
 - ◆ Kernel requests memory for structures of varying sizes
 - ◆ Some kernel memory needs to be contiguous







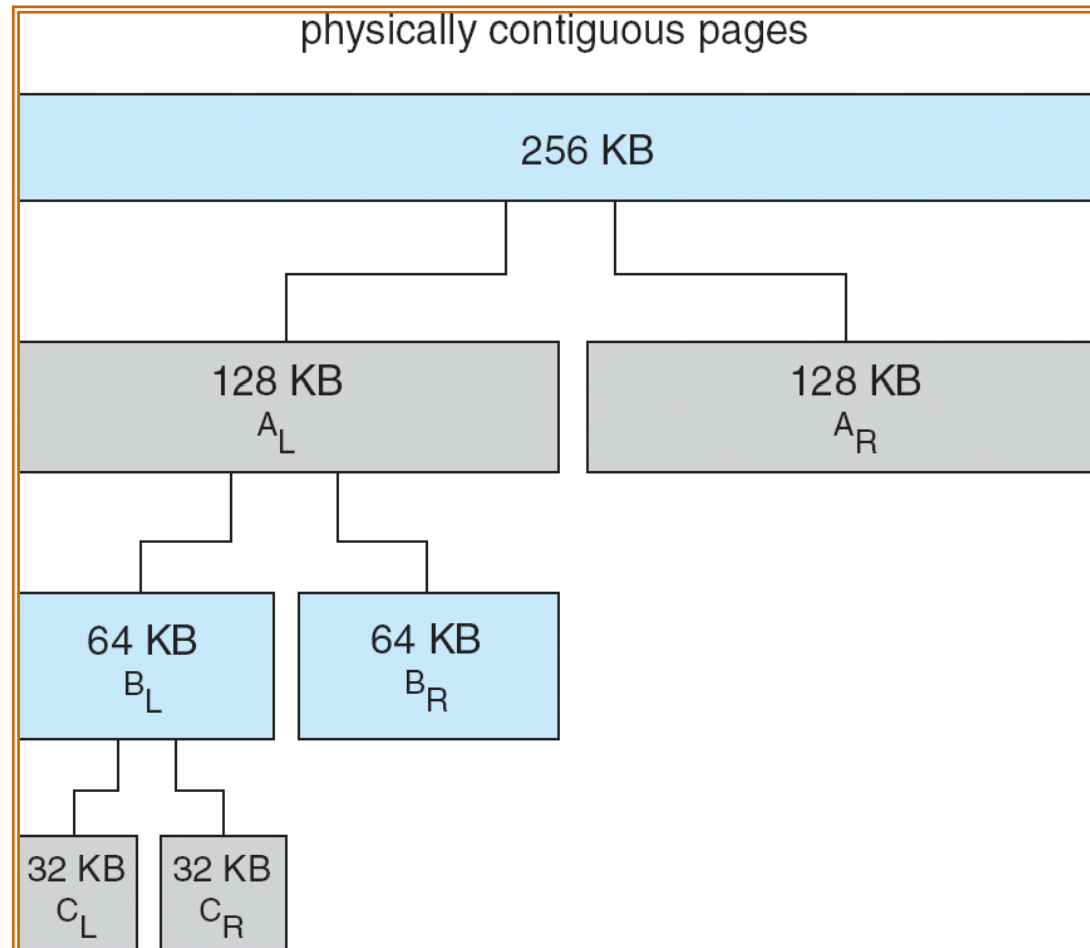
Buddy System


- Allocates memory from fixed-size segment consisting of physically-contiguous pages
- Memory allocated using **power-of-2 allocator**
 - ◆ Satisfies requests in units sized as power of 2
 - ◆ Request rounded up to next highest power of 2
 - ◆ When smaller allocation needed than is available, current chunk split into two buddies of next-lower power of 2
 - ✓ Continue until appropriate sized chunk available






Buddy System Allocator






Slab Allocator


- Alternate strategy
- **Slab** is one or more physically contiguous pages
- **Cache** consists of one or more slabs
- Single cache for each unique kernel data structure
 - ◆ Each cache filled with **objects** – instantiations of the data structure








Slab Allocator


- When cache created, filled with objects marked as **free**
- When structures stored, objects marked as **used**
- If slab is full of used objects, next object allocated from empty slab
 - ◆ If no empty slabs, new slab allocated
- Benefits include no fragmentation, fast memory request satisfaction






Illustrate the Slab Allocation






Chapter 9: Virtual Memory

- Background
- Demand Paging
- Copy-on-Write
- Page Replacement
- Allocation of Frames
- Thrashing
- Memory-Mapped Files
- Allocating Kernel Memory
- Other Considerations
- Operating-System Examples







Other Issues -- Prepaging

■ Prepaging


- ◆ To reduce the large number of page faults that occurs at process startup
- ◆ Prepage all or some of the pages a process will need, before they are referenced
- ◆ But if prepaged pages are unused, I/O and memory was wasted
- ◆ Assume s pages are prepaged and α of the pages is used
 - ✓ Is cost of $s * \alpha$ save pages faults $>$ or $<$ than the cost of prepaging
 $s * (1 - \alpha)$ unnecessary pages?
 - ✓ α near zero \Rightarrow prepaging loses






Other Issues – Page Size


- Page size selection must take into consideration:
 - ◆ fragmentation
 - ◆ table size
 - ◆ I/O overhead
 - ◆ locality






Other Issues – TLB Reach

- TLB Reach - The amount of memory accessible from the TLB
- $\text{TLB Reach} = (\text{TLB Size}) \times (\text{Page Size})$
- Ideally, the working set of each process is stored in the TLB
 - ◆ Otherwise there is a high degree of page faults
- Increase the Page Size
- Provide Multiple Page Sizes
 - ◆ This allows applications that require larger page sizes the opportunity to use them without an increase in fragmentation





Other Issues – Program Structure

■ Program structure

◆ **int A[][] = new int[1024][1024];**


◆ Each row is stored in one page


◆ Program 1 **for (j = 0; j < A.length; j++)**
 for (i = 0; i < A.length; i++)
 A[i,j] = 0;

1024 x 1024 page faults

◆ Program 2 **for (i = 0; i < A.length; i++)**
 for (j = 0; j < A.length; j++)
 A[i,j] = 0;


◆ 1024 page faults








Other Issues – I/O interlock


- **I/O Interlock** – Pages must sometimes be locked into memory
- Consider I/O - Pages that are used for copying a file from a device must be locked from being selected for eviction by a page replacement algorithm






Reason Why Frames Used For I/O Must Be In Memory






Chapter 9: Virtual Memory

- Background
- Demand Paging
- Copy-on-Write
- Page Replacement
- Allocation of Frames
- Thrashing
- Memory-Mapped Files
- Allocating Kernel Memory
- Other Considerations
- **Operating-System Examples**







Operating System Examples

- Windows XP


- Solaris






Windows XP

- Uses demand paging with **clustering**.
Clustering brings in pages surrounding the faulting page.
- Processes are assigned **working set minimum** and **working set maximum**
- Working set minimum is the minimum number of pages the process is guaranteed to have in memory






Windows XP (Cont.)


- A process may be assigned as many pages up to its working set maximum
- When the amount of free memory in the system falls below a threshold, **automatic working set trimming** is performed to restore the amount of free memory
- Working set trimming removes pages from processes that have pages in excess of their working set minimum






Solaris


- Maintains a list of free pages to assign faulting processes
- *Lotsfree* – threshold parameter (amount of free memory) to begin paging
- *Desfree* – threshold parameter to increasing paging
- *Minfree* – threshold parameter to being swapping






Solaris (Cont.)

- Paging is performed by *pageout* process
- Pageout scans pages using modified clock algorithm
- *Scanrate* is the rate at which pages are scanned. This ranges from *slowscan* to *fastscan*
- Pageout is called more frequently depending upon the amount of free memory available





Solaris 2 Page Scanner

